


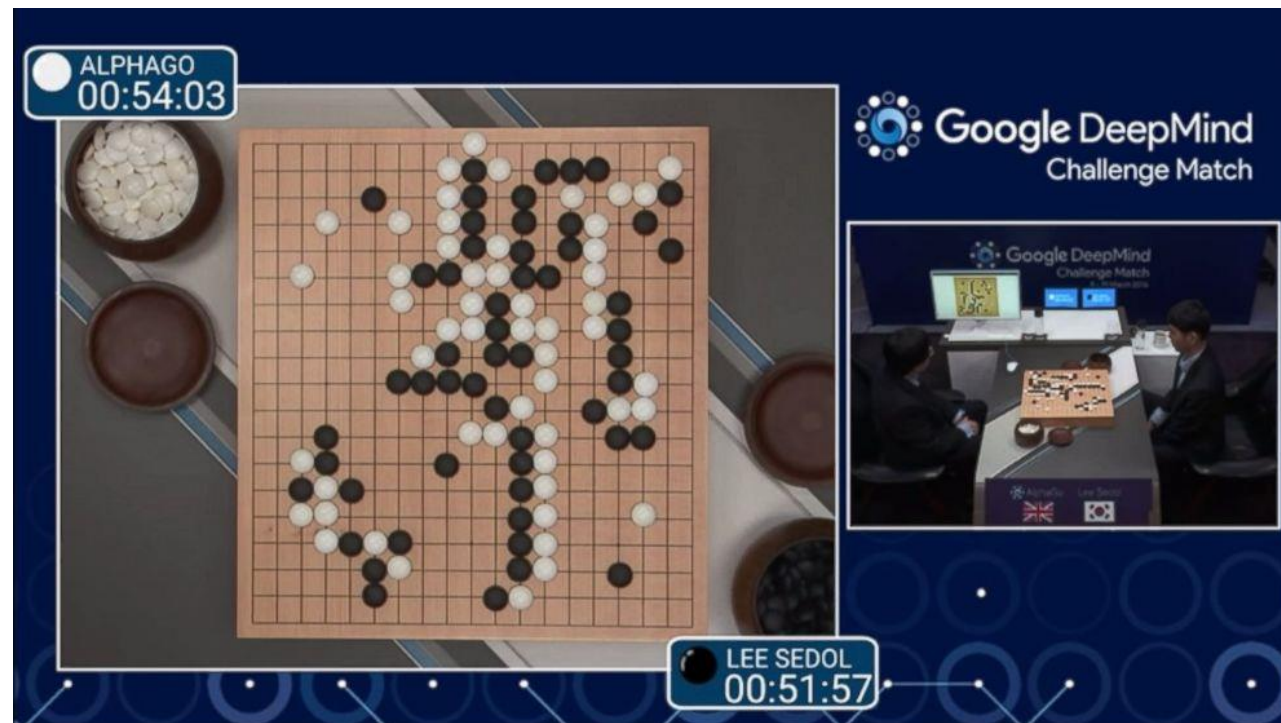
Search in Games

Chen-Yu Wei



Bernstein

Computer



Turn-Based Two-Player Game

You choose one of the three bins. I choose a number from that bin. Your goal is to maximize the chosen number.

A	
-50	50

B	
1	3

C	
15	-5


If I am

- adversarial
- random
- benign/cooperative

Turn-Based Two-Player Zero-Sum Games

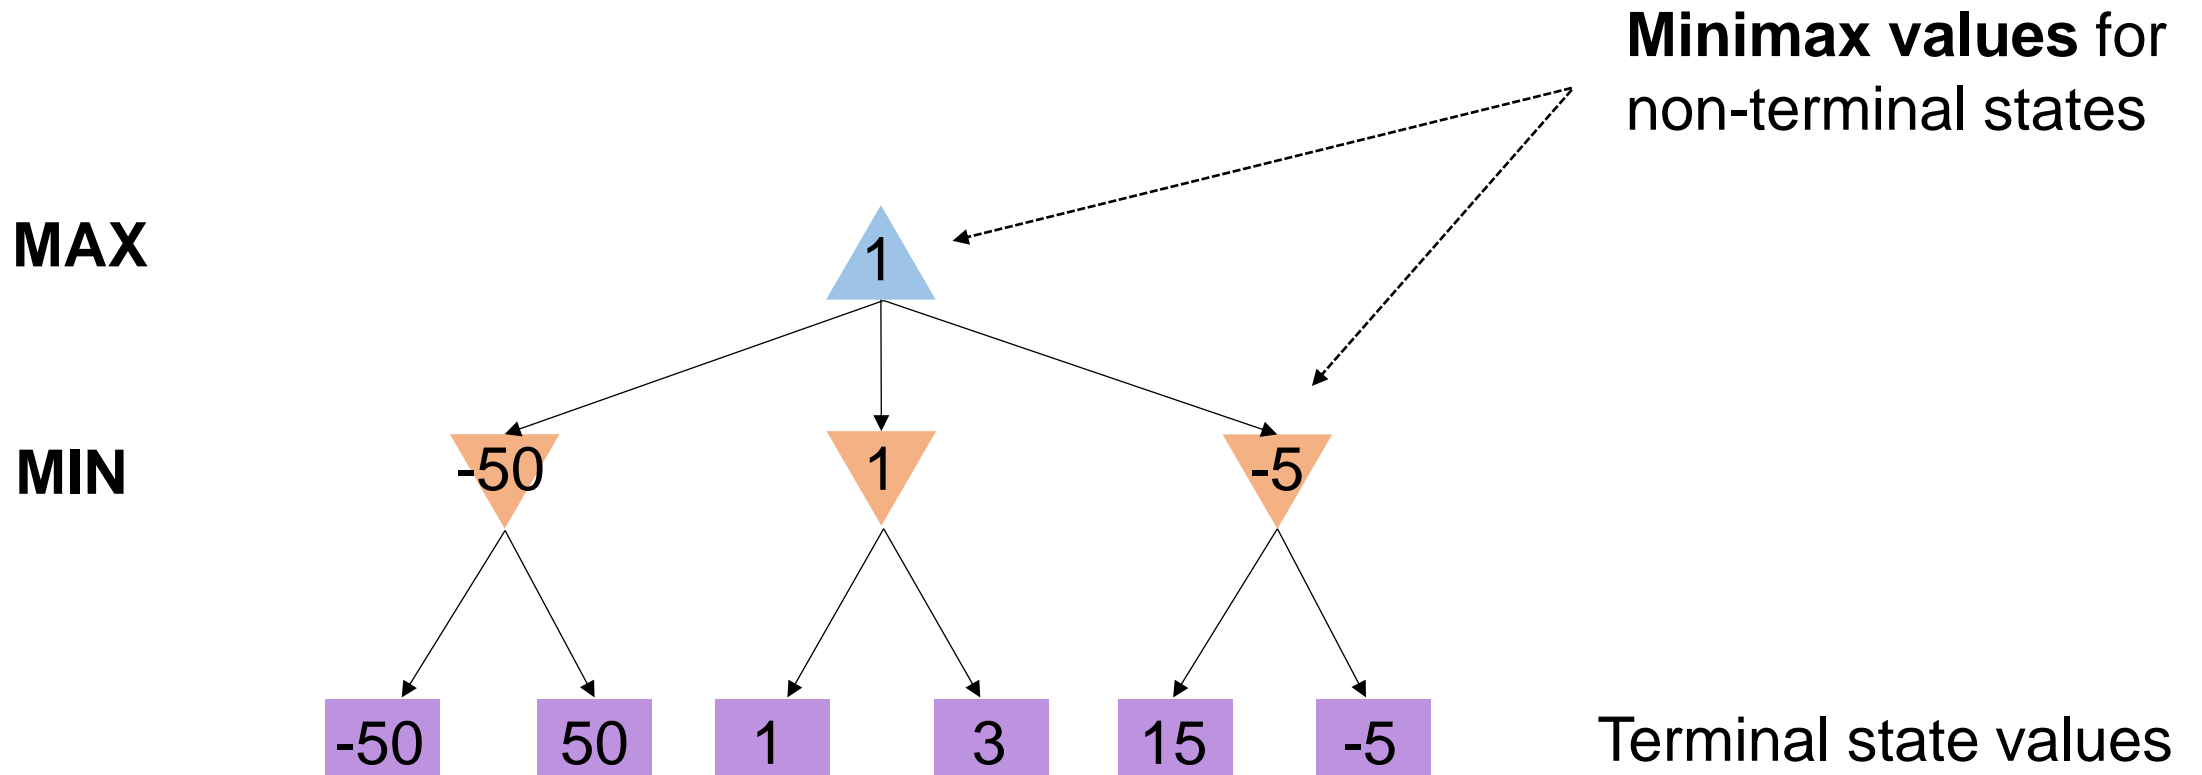
MAX

MIN




Terminal state values


Turn-Based Two-Player Zero-Sum Games



Example: PACMAN



Example: Tic-Tac-Toe



Calculating Minimax Values

```
def value(state):
```

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is MIN: return min-value(state)

```
def max-value(state):
```

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return v

```
def min-value(state):
```

initialize $v = +\infty$

for each successor of state:

$v = \min(v, \text{value}(\text{successor}))$

return v

The Minimax Policy

“**Policy**” is mapping from state to action.


“**Minimax policy**” is the optimal policy against the most adversarial opponent.

→ **MAX Player's** minimax policy

→ **MIN Player's** minimax policy

MAX

MIN




Time / Space Complexity

- Same as DFS
 - Time: $O(b^m)$
 - Space: $O(bm)$
- For chess
 - $b \approx 35$, $m \approx 100$
 - Too large to find the true minimax value/policy

Alpha-Beta Pruning and Evaluation Functions

MAX

MIN




MAX

MIN

MAX

MIN



MAX


MIN

MAX

MIN

≤ 8

8




MAX

MIN

MAX

MIN



MAX

MIN

MAX


MIN

8

7

$=$

≥ 7




MAX

MIN

MAX


MIN



MAX

MIN


MAX

MIN

MAX

MIN

MAX


MIN

MAX

MIN

MAX

MIN




MAX

MIN

MAX

MIN




MAX

MIN

MAX


MIN



MAX

MIN

MAX


MIN

MAX

MIN

MAX

MIN




MAX

MIN

MAX

MIN




MAX

MIN

MAX

MIN




MAX

MIN

MAX

MIN




MAX

MIN

MAX

MIN




MAX

MIN

MAX

MIN




MAX

MIN

MAX

MIN




MAX

MIN

MAX

MIN




MAX


MIN

MAX

MIN



Alpha-Beta Pruning



Alpha-Beta Pruning

α : MAX's best option on path to root
 β : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$  return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```


```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$  return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```

Alpha-Beta Pruning

- The pruning has **no effect** on the minimax value computed for the root.
- Child ordering affects the efficiency
 - If a MAX node finds a larger children value (or a MIN node finds a smaller children value) quicker, then more time can be saved.
- With perfect ordering, the time complexity drops to $O(b^{m/2})$
 - Doubles solvable depth
 - Full search of, e.g., chess, is still hopeless


Resource Limits

- In realistic games, cannot search to leaves
- **Solution:** depth-limited search
 - Search only to a limited depth
 - At the last layer of the search, call the **evaluation function** (heuristic function)
- **Example**
 - Suppose we have 100 seconds, can explore 10K nodes / sec
 - So can check 1M nodes per move
 - α - β reaches about depth 8 – decent chess program
- Use iterative deepening for an anytime algorithm



Evaluation Functions

- Evaluation functions score non-terminal nodes in depth-limited search



- e.g., weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

where $f_1(s) = (\text{num white queens} - \text{num black queens})$, etc.

Next Lecture

- Expectimax
- Monte-Carlo Tree Search