

Linear Contextual Bandits

Chen-Yu Wei


Contextual Bandits



all-user recommendation system



Context



personalized recommendation system

e.g. the user's historical
purchase record, location,
social network activity, ...

Contextual Bandits

For time $t = 1, 2, \dots, T$:

Environment generates a context $x_t \in \mathcal{X}$

Learner chooses an action $a_t \in \mathcal{A}$

Learner observes $r_t = R(x_t, a_t) + w_t$


$$\begin{aligned} \text{Regret} &= \max_{\pi} \sum_{t=1}^T R(x_t, \pi(x_t)) - \sum_{t=1}^T R(x_t, a_t) & \text{Optimal policy: } \pi(x) &= \operatorname{argmax}_{a \in \mathcal{A}} R(x, a) \\ &= \sum_{t=1}^T \max_{a \in \mathcal{A}} R(x_t, a) - \sum_{t=1}^T R(x_t, a_t) \end{aligned}$$

View Each Context as a Separate MAB

$$\begin{aligned}\text{Regret} &= \sum_{t=1}^T \max_{a \in \mathcal{A}} R(x_t, a) - \sum_{t=1}^T R(x_t, a_t) \\ &= \sum_{x \in \mathcal{X}} \left(\sum_{t: x_t = x} \max_{a \in \mathcal{A}} R(x, a) - \sum_{t: x_t = x} R(x, a_t) \right)\end{aligned}$$


Not scalable and not generalizable

Function Approximation in Contextual Bandits




Find an f so that $f(x) \approx y$ for **seen** (x, y) pairs
Hoping that $f(x') \approx y'$ also holds for **unseen** x'

x : context, a : action, r : reward



value-based approach



policy-based approach

If a good approximation \hat{R} is found, a good policy can be derived as

$$\pi(x) = \operatorname{argmax}_a \hat{R}(x, a)$$

Linear Contextual Bandits

This is a linear **assumption**, not just linear **function approximation**. The former is stronger.

Linear Reward Assumption: $R(x, a) = \phi(x, a)^\top \theta^*$

$\phi(x, a) \in \mathbb{R}^d$ is a **feature vector** for the context-action pair (known to learner)

$\theta^* \in \mathbb{R}^d$ is the ground-truth **weight vector** (hidden from learner)

Given: feature mapping $\phi: \mathcal{X} \times \mathcal{A} \rightarrow \mathbb{R}^d$

For time $t = 1, 2, \dots, T$:

Environment generates a context $x_t \in \mathcal{X}$

Learner chooses an action $a_t \in \mathcal{A}$

Learner observes $r_t = \phi(x_t, a_t)^\top \theta^* + w_t$ (w_t is zero-mean)

$$\text{Regret} = \sum_{t=1}^T \max_{a \in \mathcal{A}} R(x_t, a) - \sum_{t=1}^T R(x_t, a_t) = \sum_{t=1}^T \max_{a \in \mathcal{A}} \phi(x_t, a)^\top \theta^* - \sum_{t=1}^T \phi(x_t, a_t)^\top \theta^*$$

Linear CB is a Generalization of MAB

Key Questions in Linear Contextual Bandits

- How to obtain an estimated reward function $\hat{R}(x, a)$?
 - Was easy in multi-armed bandits – today we'll see how to do this in linear CB
- How to explore?
 - ϵ -greedy

$$a_t = \begin{cases} \text{uniform}(\mathcal{A}) & \text{with prob. } \epsilon \\ \operatorname{argmax}_a \hat{R}_t(x_t, a) & \text{with prob. } 1 - \epsilon \end{cases}$$

- Boltzmann exploration

$$p_t(a) \propto \exp(\lambda_t \hat{R}_t(x_t, a))$$

- Optimism in the face of uncertainty (LinUCB)
- Thompson Sampling

How to Estimate the Reward Function $R(x, a)$?

- Recall $R(x, a) = \phi(s, a)^\top \theta^*$. We only need to estimate θ^* .
- At time t , we already gathered

$$r_1 = \phi(x_1, a_1)^\top \theta^* + w_1$$

$$r_2 = \phi(x_2, a_2)^\top \theta^* + w_2$$

\vdots

$$r_{t-1} = \phi(x_{t-1}, a_{t-1})^\top \theta^* + w_{t-1}$$

How to estimate θ^* ?

Linear Regression

Linear Regression

At time t , we have collected $(x_1, a_1, r_1), (x_2, a_2, r_2), \dots, (x_{t-1}, a_{t-1}, r_{t-1})$.

We want to generate an estimation $\hat{\theta}_t$ such that $\phi(x_i, a_i)^\top \hat{\theta}_t \approx r_i$

Linear Regression / Ridge Regression (define $\phi_i = \phi(x_i, a_i)$)

$$\hat{\theta}_t = \min_{\theta} \sum_{i=1}^{t-1} (\phi_i^\top \theta - r_i)^2 + \lambda \|\theta\|^2 \quad \Leftrightarrow \quad \hat{\theta}_t = \left(\lambda I + \sum_{i=1}^{t-1} \phi_i \phi_i^\top \right)^{-1} \left(\sum_{i=1}^{t-1} \phi_i r_i \right)$$

$\Rightarrow \hat{R}_t(x, a) = \phi(x, a)^\top \hat{\theta}_t$ (Use this directly in ϵ -greedy or Boltzmann exploration!)

To design a UCB algorithm, we have to quantify the estimation error $\hat{\theta}_t - \theta^*$

What can we say about $\hat{\theta}_t - \theta^*$?

Let's develop some intuition first.. (This intuition comes from Haipeng Luo's [lecture](#))

Let $r_i = \phi_i^\top \theta^* + w_i$ for $i = 1, \dots, N$

Assume $w_i \sim \mathcal{N}(0, \sigma^2)$, and

Assume $\{\phi_1, \dots, \phi_N\}$ are fixed vectors independent from $\{w_1, \dots, w_N\}$

Let

$$\hat{\theta} = \left(\sum_{i=1}^N \phi_i \phi_i^\top \right)^{-1} \left(\sum_{i=1}^N \phi_i r_i \right)$$

Question: What can we say about $\hat{\theta} - \theta^*$?

Geometric Intuition

Concentration Inequality for Linear Regression

Theorem.

In linear contextual bandits, assume w_t is zero-mean and 1-sub-Gaussian.
 $\|\phi(x, a)\|_2 \leq 1$, $\|\theta^*\|_2 \leq 1$.

Let

$$\hat{\theta}_t = \Lambda_t^{-1} \left(\sum_{i=1}^{t-1} \phi_i r_i \right), \quad \text{where } \Lambda_t = I + \sum_{i=1}^{t-1} \phi_i \phi_i^\top.$$

Then with probability at least $1 - \delta$, for all $t = 1, \dots, T$,

$$\|\theta^* - \hat{\theta}_t\|_{\Lambda_t}^2 \leq \beta \triangleq d \log \left(1 + \frac{T}{d} \right) + 3 \log \frac{1}{\delta}$$

Another Viewpoint on the Concentration Inequality

$$\begin{aligned}\|\theta^* - \hat{\theta}_t\|_{\Lambda_t}^2 &= (\theta^* - \hat{\theta}_t)^\top \left(I + \sum_{i=1}^{t-1} \phi_i \phi_i^\top \right) (\theta^* - \hat{\theta}_t) \\ &= \underbrace{\sum_{i=1}^{t-1} (\phi_i^\top \theta^* - \phi_i^\top \hat{\theta}_t)^2}_{\text{The difference between the predictions of } \theta^* \text{ and } \hat{\theta}_t \text{ over the past samples}} + \|\theta^* - \hat{\theta}_t\|^2 = O(d \log(T/\delta))\end{aligned}$$

The difference between the predictions of θ^* and $\hat{\theta}_t$ over the past samples

LinUCB

Most “optimistic” estimation for the reward of a

LinUCB

In round t , receive x_t , draw

$$a_t = \operatorname{argmax}_{a \in \mathcal{A}}$$

Observe $r_t = \phi(x_t, a_t)^\top \theta^* + w_t$.

LinUCB

LinUCB

In round t , receive x_t , draw

$$a_t = \operatorname{argmax}_{a \in \mathcal{A}} \max_{\theta: \|\theta - \hat{\theta}_t\|_{\Lambda_t} \leq \beta} \phi(x_t, a)^\top \theta$$

where

$$\hat{\theta}_t = \Lambda_t^{-1} \left(\sum_{i=1}^{t-1} \phi_i r_i \right), \quad \Lambda_t = I + \sum_{i=1}^{t-1} \phi_i \phi_i^\top.$$

Observe $r_t = \phi(x_t, a_t)^\top \theta^* + w_t$.

LinUCB

LinUCB

In round t , receive x_t , draw

$$a_t = \operatorname{argmax}_{a \in \mathcal{A}} \quad \phi(x_t, a)^\top \hat{\theta}_t + \sqrt{\beta} \|\phi(x_t, a)\|_{\Lambda_t^{-1}}$$

where

$$\hat{\theta}_t = \Lambda_t^{-1} \left(\sum_{i=1}^{t-1} \phi_i r_i \right), \quad \Lambda_t = I + \sum_{i=1}^{t-1} \phi_i \phi_i^\top.$$

Observe $r_t = \phi(x_t, a_t)^\top \theta^* + w_t$.

Regret Analysis for LinUCB

Regret Bound of LinUCB

With probability at least $1 - \delta$,

$$\text{Regret} \leq O(d\sqrt{T} \log(T/\delta)) = \tilde{O}(d\sqrt{T}) .$$

Elliptical Potential Lemma

Let $\phi_i \in \mathbb{R}^d$ and $\|\phi_i\|_2 \leq 1$. Define $\Lambda_t = I + \sum_{i=1}^{t-1} \phi_i \phi_i^\top$.

Then

$$\sum_{t=1}^T \|\phi_t\|_{\Lambda_t^{-1}}^2 \leq d \log \left(1 + \frac{T}{d} \right).$$

Thompson Sampling

Thompson Sampling for Linear Contextual Bandits

In round t , receive x_t , draw

$$\theta_t \sim \mathcal{N}(\hat{\theta}_t, \Lambda_t^{-1})$$

$$a_t = \operatorname{argmax}_{a \in \mathcal{A}} \phi(x_t, a)^\top \theta_t$$

where

$$\hat{\theta}_t = \Lambda_t^{-1} \left(\sum_{i=1}^{t-1} \phi_i r_i \right), \quad \Lambda_t = I + \sum_{i=1}^{t-1} \phi_i \phi_i^\top.$$

Observe $r_t = \phi(x_t, a_t)^\top \theta^* + w_t$.

There is no assumption on the distribution of x_t

- How is this possible?