CHEN-YU WEI

Homepage | Google Scholar | E-mail: chenyu.wei@virginia.edu | Last update: March 2024

EDUCATION

University of Southern California — Los Angeles, CA
Ph.D. in Computer Science — 2017–2022
Supervisor: Haipeng Luo
Thesis: Robust and Adaptive Online Decision Making

National Taiwan University — Taipei, Taiwan
M.S. in Communication Engineering — 2013–2015
Supervisor: Wanjiun Liao
Thesis: Downlink Scheduling Policies in Heterogeneous Networks with User Equipment Side Interference Cancellation

National Taiwan University — Taipei, Taiwan
B.S. in Electrical Engineering — 2008–2012

EXPERIENCES

University of Virginia — Charlottesville, VA
Assistant Professor — 2023–Present
Computer Science Department

MIT Institute for Data, Systems, and Society (IDSS) — Cambridge, MA
Postdoctoral Associate — Spring and Summer 2023
Supervisor: Alexander Rakhlin

Simons Institute — Berkeley, CA
Research Fellow — Fall 2022
Program: Data-Driven Decision Processes

Simons Institute — Berkeley, CA
Student Visitor — Spring 2022
Program: Learning and Games

Google Research — Remote
Research Intern — Summer 2021
Supervisor: Christoph Dann, Julian Zimmert
Topic: Corruption Robust Reinforcement Learning

Simons Institute — Remote
Student Visitor — Fall 2020
Program: Theory of Reinforcement Learning

Microsoft Research — Redmond, WA
Research Intern — Summer 2020
Supervisor: Alekh Agarwal
Topic: Personalized Federated Learning

Yahoo Research — New York City
Research Intern — Summer 2019
Supervisor: Alina Beygelzimer
Topic: Bandit Classification

Academia Sinica — Taipei, Taiwan
Research Intern — Spring 2012
Supervisor: Yi-Hsuan Yang
Topic: Music Information Retrieval

Stanford University — Palo Alto, CA
Research Intern (Undergraduate Visiting Research (UGVR) Program) — Summer 2011
Supervisor: Boris Murmann
Topic: Circuit Design for Medical Ultrasound

HONORS AND AWARDS

Finalist (Top 2 in CS), Best Dissertation Award, USC Viterbi Engineering School	*2023*
Top Reviewers, NeurIPS	*2022*
Prize for Excellence in Research with a Substantial Mathematical Component, Center for Applied Math Science, USC	*2022*
Simons-Berkeley Research Fellowship, Simons Institute for the Theory of Computing	*2022*
Best Paper Award, International Conference on Algorithmic Learning Theory	*2022*
Best Paper Award, Conference on Learning Theory	*2021*
Best Research Assistant Award, Computer Science Department, USC	*2020*
Best Poster Award, SoCal Machine Learning Symposium	*2019*
Taiwan-USC Scholarship, Ministry of Education, Taiwan	*2017*
Tenth Place, ACM International Collegiate Programming Contest – Asia Regional	*2010*

PUBLICATIONS (CONFERENCE PAPERS)

Near-Optimal Policy Optimization for Correlated Equilibrium in General-Sum Markov Games — *AISTAT 2024*
(α-β) Yang Cai, Haipeng Luo, Chen-Yu Wei, Weiqiang Zheng (Oral)

Towards Optimal Regret in Linear MDPs with Bandit Feedback — *ICLR 2024*
(α-β) Haolin Liu, Chen-Yu Wei, Julian Zimmert (Spotlight)

Bypassing the Simulator: Near-Optimal Adversarial Linear Contextual Bandits — *NeurIPS 2023*
(α-β) Haolin Liu, Chen-Yu Wei, Julian Zimmert

Last-Iterate Convergent Policy Gradient Primal-Dual Methods for Constrained MDPs — *NeurIPS 2023*
Dongsheng Ding*, Chen-Yu Wei*, Kaiqing Zhang*, Alejandro Ribeiro

No-Regret Online Reinforcement Learning with Adversarial Losses and Transitions — *NeurIPS 2023*
Tiancheng Jin*, Junyan Liu*, Chloe Rouyer, William Chang, Chen-Yu Wei, Haipeng Luo

First- and Second-Order Bounds for Adversarial Linear Contextual Bandits — *NeurIPS 2023*
Julia Olkhovskaya, Jack Mayo, Tim van Erven, Gergely Neu, Chen-Yu Wei

Uncoupled and Convergent Learning in Two-Player Zero-Sum Markov Games — *NeurIPS 2023*
(α-β) Yang Cai, Haipeng Luo, Chen-Yu We*, Weiqiang Zheng

A Blackbox Approach to Best of Both Worlds in Bandits and Beyond — *COLT 2023*
(α-β) Christoph Dann, Chen-Yu Wei, Julian Zimmert

Best of Both Worlds Policy Optimization — *ICML 2023*
(α-β) Christoph Dann, Chen-Yu Wei, Julian Zimmert (Long talk)

Refined Regret for Adversarial MDPs with Linear Function Approximation — *ICML 2023*
(α-β) Yan Dai, Haipeng Luo, Chen-Yu Wei, Julian Zimmert

A Unified Algorithm for Stochastic Path Problems — *ALT 2023*
(α-β) Christoph Dann, Chen-Yu Wei, Julian Zimmert

Independent Policy Gradient for Large-Scale Markov Potential Games: Sharper Rates, Function Approximation, and Game-Agnostic Convergence — *ICML 2022*
Dongsheng Ding*, Chen-Yu Wei*, Kaiqing Zhang*, Mihailo Jovanovic (Long talk)

Personalization Improves Privacy-Accuracy Tradeoffs in Federated Optimization — *ICML 2022*
Alberto Bietti, Chen-Yu Wei, Miroslav Dudik, John Langford, Zhiwei Steven Wu

A Model Selection Approach for Corruption Robust Reinforcement Learning — *ALT 2022*
Chen-Yu Wei, Christoph Dann, Julian Zimmert (Best Paper Award)

Decentralized Cooperative Reinforcement Learning with Hierarchical Information Structure — *ALT 2022*
Hsu Kao, Chen-Yu Wei, Vijay Subramanian

Policy Optimization in Adversarial MDPs: Improved Exploration via Dilated Bonuses — *NeurIPS 2021*
Haipeng Luo*, Chen-Yu Wei*, Chung-Wei Lee

Achieving Near Instance-Optimality and Minimax-Optimality in Stochastic and Adversarial Linear Bandits Simultaneously — *ICML 2021*
(α-β) Chung-Wei Lee, Haipeng Luo, Chen-Yu Wei, Mengxiao Zhang, Xiaojin Zhang

Non-stationary RL without Prior Knowledge: An Optimal Black-box Approach — *COLT 2021*
Chen-Yu Wei, Haipeng Luo (Best Paper Award)

Last-iterate Convergence of Decentralized Optimistic Gradient Descent/Ascent in Infinite-horizon Competitive Markov Games Chen-Yu Wei, Chung-Wei Lee*, Mengxiao Zhang*, Haipeng Luo	*COLT 2021*
Impossible Tuning Made Possible: A New Expert Algorithm and Its Applications (α-β) Liyu Chen, Haipeng Luo, Chen-Yu Wei	*COLT 2021*
Minimax Regret for Stochastic Shortest Path with Adversarial Costs and Known Transition Liyu Chen, Haipeng Luo, Chen-Yu Wei	*COLT 2021*
Learning Infinite-horizon Average-reward MDPs with Linear Function Approximation Chen-Yu Wei, Mehdi Jafarnia-Jahromi, Haipeng Luo, Rahul Jain	*AISTAT 2021*
Linear Last-iterate Convergence for Constrained Saddle-point Optimization Chen-Yu Wei, Chung-Wei Lee, Mengxiao Zhang, Haipeng Luo	*ICLR 2021*
Adversarial Online Learning with Changing Action Sets: Efficient Algorithms with Approximate Regret Bounds Ehsan Emamjomeh-Zadeh*, Chen-Yu Wei*, Haipeng Luo, David Kempe	*ALT 2021*
Bias No More: High-probability Data-dependent Regret Bounds for Adversarial Bandits and MDPs (α-β) Chung-Wei Lee, Haipeng Luo, Chen-Yu Wei, Mengxiao Zhang (Oral)	*NeurIPS 2020*
Taking a Hint: How to Leverage Loss Predictors in Contextual Bandits? Chen-Yu Wei, Haipeng Luo, Alekh Agarwal	*COLT 2020*
Model-free Reinforcement Learning in Infinite-horizon Average-reward Markov Decision Processes Chen-Yu Wei, Mehdi Jafarnia-Jahromi, Haipeng Luo, Hiteshi Sharma, Rahul Jain	*ICML 2020*
A New Algorithm for Non-stationary Contextual Bandits: Efficient, Optimal, and Parameter-free (α-β) Yifang Chen, Chung-Wei Lee, Haipeng Luo, Chen-Yu Wei	*COLT 2019*
Improved Path-length Regret Bounds for Bandits (α-β) Sébastien Bubeck, Yuanzhi Li, Haipeng Luo, Chen-Yu Wei	*COLT 2019*
Bandit Multiclass Linear Classification: Efficient Algorithms for the Separable Case (α-β) Alina Beygelzimer, Dávid Pál, Balázs Szörényi, Devanathan Thiruvenkatachari, Chen-Yu Wei, Chicheng Zhang	*ICML 2019*
Beating Stochastic and Adversarial Semi-bandits Optimally and Simultaneously Julian Zimmert, Haipeng Luo, Chen-Yu Wei (Long talk)	*ICML 2019*
Efficient Online Portfolio with Logarithmic Regret (α-β) Haipeng Luo, Chen-Yu Wei, Kai Zheng (Spotlight)	*NeurIPS 2018*
More Adaptive Algorithms for Adversarial Bandits Chen-Yu Wei, Haipeng Luo	*COLT 2018*
Efficient Contextual Bandits in Non-stationary Worlds Haipeng Luo*, Chen-Yu Wei*, Alekh Agarwal, John Langford	*COLT 2018*
Online Reinforcement Learning in Stochastic Games Chen-Yu Wei, Yi-Te Hong, Chi-Jen Lu	*NeurIPS 2017*
Tracking the Best Expert in Non-stationary Stochastic Environments Chen-Yu Wei, Yi-Te Hong, Chi-Jen Lu	*NeurIPS 2016*

PUBLICATIONS (WORKSHOP PAPERS)

Federated Residual Learning *NeurIPS Workshop on Scalability, Privacy, and Security in Federated Learning (Spicy-FL)* Chen-Yu Wei, Alekh Agarwal, John Langford	*2020*
Analyzing the Variance of Policy Gradient Estimators for the Linear-Quadratic Regulator *NeurIPS Workshop on Optimization Foundations for Reinforcement Learning (OPTRL)* Sébastien Arnold*, James Preiss*, Chen-Yu Wei*, Marius Kloft	*2019*
Understanding the Variance of Policy Gradient Estimators in Reinforcement Learning *SoCal Machine Learning Symposium (SoCalML)* Sébastien Arnold*, James Preiss*, Chen-Yu Wei*, Marius Kloft (Best Poster Award)	*2019*

INVITED TALKS

Exploration Bonus for Policy Optimization, AI/ML Seminar, UVa CS	*Sep. 2023*
Exploration Bonus for Policy Optimization, Distinguished Talk Series, Microsoft Research	*Jan. 2023*
Some Recent Advances in the Theory of Online Decision Making, Special Topics, National Taiwan University	*Oct. 2022*
Optimal Dynamic Regret for Bandits without Prior Knowledge, BLISS Seminar, UC Berkeley	*Oct. 2022*
Optimal Dynamic Regret for Bandits without Prior Knowledge, D3P program workshop, Simons Institute	*Sep. 2022*
Robust and Adaptive Online Decision Making, UMich ECE Seminar	*Apr. 2022*
Robust and Adaptive Online Decision Making, UVa CS Seminar	*Mar. 2022*
Non-stationary RL without Prior Knowledge: an Optimal Black-box Approach, COLT Best Paper Talk	*Aug. 2021*
Linear Last-iterate Convergence of Constrained Saddle-point Optimization, UW Learning in Games Seminar	*May. 2021*
Learning Infinite-horizon Average-reward MDPs with Linear Function Approximation, RL Virtual Seminars	*Sep. 2020*
Bandit Multiclass Linear Classification: Efficient Algorithms for the Separable Case, Theory Day, UC Riverside	*Jan. 2020*
Bandit Multiclass Linear Classification: Efficient Algorithms for the Separable Case, Theory Lunch, MSR	*June 2019*
Beating Stochastic and Adversarial Semi-bandits Optimally and Simultaneously, ICML Long Talk	*June 2019*
Efficient Online Portfolio with Logarithmic Regret, NeurIPS Spotlight Talk	*Dec. 2018*

ACADEMIC ACTIVITIES

UVA CS 4501: Introduction to Algorithmic Economics
Instructor   *Spring 2024*
Co-teaching with Denis Nekipelov

UVA CS 6501: Reinforcement Learning
Instructor   *Spring 2024*

USC CSCI 567: Machine Learning
Teaching Assistant   *Fall 2021*
Instructor: Haipeng Luo

USC CSCI 270: Introduction to Algorithms and Theory of Computing course
Teaching Assistant   *Spring 2021*
Instructor: Shawn Shamsian

USC CSCI 699: Introduction to Online Learning
Teaching Assistant   *Fall 2017*
Instructor: Haipeng Luo

Reviewer for COLT, ALT, STOC, FOCS, NeurIPS, ICML, ICLR, AISTAT, AAAI, JMLR, MOR, TMLR