

Policy Evaluation

Chen-Yu Wei

Policy Evaluation

Given: a policy π

Evaluate $V^\pi(s)$ or $Q^\pi(s, a)$

On-policy policy evaluation: the learner can execute π to evaluate π

$\pi_b(\cdot|s)$

Off-policy/offline policy evaluation: the learner can only execute some $\pi_b \neq \pi$, or can only access some existing dataset to evaluate π

(s, a, r, s')

$(s_1, a_1, r_1, s_2, a_2, \dots)$

↑
behavior policy

Use cases:

- Approximate policy iteration: $\pi^{(k)}(s) = \operatorname{argmax}_a Q^{\pi^{(k-1)}}(s, a)$
- Estimate the value of a policy before deploying it in the real world, e.g., COVID-related border measures, economic recovery policies, or policy changes in recommendation systems.

Value Iteration for V^π / Q^π

$$\boxed{Q^*(s,a)} \quad \boxed{V^*(s)}$$

\downarrow
 $\pi^*(s) = \underset{a}{\operatorname{argmax}} Q^*(s,a)$

Input: π

For $k = 1, 2, \dots$

$$\forall s, \quad V^{(k)}(s) \leftarrow \sum_a \pi(a|s) \left(R(s,a) + \gamma \sum_{s'} P(s'|s,a) V^{(k-1)}(s') \right)$$

$V^{(k)}(s) \xrightarrow{k \rightarrow \infty} V^*(s)$

Input: π

For $k = 1, 2, \dots$

$$\forall s, a, \quad Q^{(k)}(s,a) \leftarrow R(s,a) + \gamma \sum_{s',a'} P(s'|s,a) \pi(a'|s') Q^{(k-1)}(s',a')$$

$Q^{(k)} \rightarrow Q^*$

On-Policy Policy Evaluation

LSTD and TD

For $k = 1, 2, \dots$

Collecting samples $\{(s_i, r_i, s'_i)\}_{i=1}^n$ using π

$$\theta_k \leftarrow \underset{\theta}{\operatorname{argmin}} \sum_{i=1}^n \left(V_{\theta}(s_i) - r_i - \gamma V_{\theta_{k-1}}(s'_i) \right)^2$$

$$a_i \sim \pi(\cdot | s_i), \quad \mathbb{E}[r_i] = R(s_i, a_i), \quad s'_i \sim p(\cdot | s_i, a_i)$$

Least-Square
Temporal Difference (LSTD)

$$V(s) \underset{\theta_k}{\approx} \sum_a \pi(a|s) \left(R(s,a) + \gamma \mathbb{E}_{s' \sim p(\cdot|s,a)} [V_{\theta_{k-1}}(s')] \right)$$

For $i = 1, 2, \dots$

Draw $a_i \sim \pi(\cdot | s_i)$

Observe reward r_i and next state s_{i+1}

$$\theta_i \leftarrow \theta_{i-1} - \alpha \nabla_{\theta} \left(V_{\theta}(s_i) - r_i - \gamma V_{\theta_{i-1}}(s_{i+1}) \right)^2$$

Temporal difference learning

TD learning

TD(0) λ

LSTDQ and TDQ \mathcal{TD}

For $k = 1, 2, \dots$

Collecting samples $\{(s_i, a_i, r_i, s'_i)\}_{i=1}^n$ using π

$$\theta_k \leftarrow \underset{\theta}{\operatorname{argmin}} \sum_{i=1}^n \left(Q_{\theta}(s_i, a_i) - r_i - \gamma \sum_{a'} \pi(a'|s'_i) Q_{\theta_{k-1}}(s'_i, a') \right)^2$$

For $i = 1, 2, \dots$

Draw $a_i \sim \pi(\cdot | s_i)$, observe reward r_i and next state s_{i+1}

$$\theta_i \leftarrow \theta_{i-1} - \alpha \nabla_{\theta} \left(Q_{\theta}(s_i, a_i) - r_i - \gamma \sum_{a'} \pi(a'|s'_i) Q_{\theta_{k-1}}(s'_i, a') \right)^2$$

TD with Linear Function Approximation

$$A \succeq 0 : A \text{ is psd}$$

$$A \succeq B \Leftrightarrow A - B \text{ is psd}$$

$$\text{BC : } R(s,a) + \gamma \mathbb{E}_{s' \sim P(\cdot|s,a)} \max_{a'} \tilde{Q}(s',a') = \phi(s,a)^T \theta^* \quad \forall \tilde{Q}$$

Let μ be the stationary state distribution under policy π . Furthermore, assume

$$(1) \quad V^\pi(s) = \phi(s)^T \theta^* \quad (\text{realizability assumption})$$

$$(2) \quad \mathbb{E}_{s \sim \mu} [\phi(s)\phi(s)^T] \succcurlyeq \rho I \text{ for some } \rho > 0 \quad (\text{coverage assumption})$$

imply (set $\tilde{Q} = Q^*$)

Then the following TD update:

Realizability in Q^* :

For $i = 1, 2, \dots$

$$Q^*(s,a) = \phi(s,a)^T \theta^*$$

In fact, even if the samples are generated as $a_i \sim \pi(\cdot|s_i)$, $r_i = R(s_i, a_i)$, $s_{i+1} \sim P(\cdot|s_i, a_i)$

Sample $s \sim \mu$, $a \sim \pi(\cdot|s)$, $r \sim R(s, a)$, $s' \sim P(\cdot|s, a)$

$$\theta_i \leftarrow \theta_{i-1} - \alpha_i (\phi(s)^T \theta_{i-1} - r - \gamma \phi(s')^T \theta_{i-1}) \phi(s)$$

converges to θ^* with properly chosen α_i .

✂

$$V^\pi(s) = \phi(s)^\top \theta^*$$

condition θ_i

$$\|\theta_{i+1} - \theta^*\|^2 = \|\theta_i - \alpha(\phi(s)^\top \theta_i - r - \gamma \phi(s')^\top \theta_i) \phi(s) - \theta^*\|^2 \text{ where } \begin{cases} s \sim \mu, & a \sim \pi(\cdot|s), & \mathbb{E}(r) = R(s,a) \\ s' \sim p(\cdot|s,a) \end{cases}$$

$$= \|\theta_i - \theta^*\|^2 - 2\alpha (\theta_i - \theta^*)^\top \underbrace{(\phi(s)^\top \theta_i - r - \gamma \phi(s')^\top \theta_i) \phi(s)}_g + \alpha^2 \|g\|^2$$

$$\mathbb{E}[\|\theta_{i+1} - \theta^*\|^2] = \|\theta_i - \theta^*\|^2 - 2\alpha (\theta_i - \theta^*)^\top \underbrace{\mathbb{E}\left[\phi(s)^\top \theta_i - r - \gamma \phi(s')^\top \theta_i\right] \phi(s)}_g + \alpha^2 \|g\|^2$$

$$= \|\theta_i - \theta^*\|^2 - 2\alpha (\theta_i - \theta^*)^\top \mathbb{E}\left[\phi(s) \left(\phi(s)^\top - \gamma \phi(s')^\top\right)\right] (\theta_i - \theta^*) + \alpha^2 \|g\|^2$$

$$\begin{aligned} \text{blue} &= \mathbb{E}\left[-V^\pi(s) + r + \gamma V^\pi(s')\right] \\ &= \mathbb{E}\left[-V^\pi(s) + \sum_a \pi(a|s) (R(s,a) + \gamma \mathbb{E}_{s' \sim p(\cdot|s,a)} (V^\pi(s')))\right] \\ &= 0 \end{aligned}$$

Comparison

Why does **Linear TDQ** converge (and converges to the correct solution) but **Linear Q-Learning** diverges?

Comparison

Under coverage assumption

(i.e., the data $\{(s_i, a_i, r_i, s_i')\}$ sufficiently cover every state-action pair / feature space)

	LSVI	Watkins's Q-Learning	On-Policy LSTD(Q) / TD(Q)
Tabular	$Q^{(k)} \rightarrow Q^*$	$Q^{(k)} \rightarrow Q^*$	$V^{(k)} \rightarrow V^\pi / Q^{(k)} \rightarrow Q^\pi$ under realizability
Linear Approx.	$Q^{(k)} \rightarrow Q^*$ under Bellman completeness	Diverges even with Bellman completeness	