

# Markov Decision Processes

Chen-Yu Wei


# Sequence of Actions



To win the game, the learner has to take a sequence of actions  $a_1 \rightarrow a_2 \rightarrow \dots \rightarrow a_H$ .  
The effect of a particular action may not be revealed instantaneously.

- Some effect may be revealed instantaneously
- Some may be revealed later


# Sequence of Actions



# Sequence of Actions

- The number of possible combinations of actions grows **exponentially** with the length of the sequence.
- We would like to decompose the problem so that every single decision in the sequence is easy to make.
- **State:** a **summary** of the **status of the world** and the **progress of the learner**, so that all future decisions can only depend on the state and not on everything else.
  - Games (Go, Chess): To decide future moves, the player only need the current board configuration.
  - Robot navigation to a goal: only need the current position and not the exact path reaching the current position.
  - Inventory management: only need the current inventory level, and not the sequence of past sales.

# Sequence of Actions



Like a sequential contextual bandit problem – except that future contexts depends on the learner's past decisions.

# Interaction Protocol (Episodic Setting)

For **episode**  $t = 1, 2, \dots, T$ :

$h \leftarrow 1$

✓ Environment generates initial state

$\chi_t$   
 $s_{t,1}$

While episode  $t$  has not ended:

Learner chooses an action  $a_{t,h}$

Learner observes instantaneous reward  $r_{t,h}$  with  $\mathbb{E}[r_{t,h}] = R(\underline{s_{t,h}}, \underline{a_{t,h}})$

Environment generates next state  $s_{t,h+1} \sim P(\cdot \mid s_{t,h}, a_{t,h})$

$h \leftarrow h + 1$

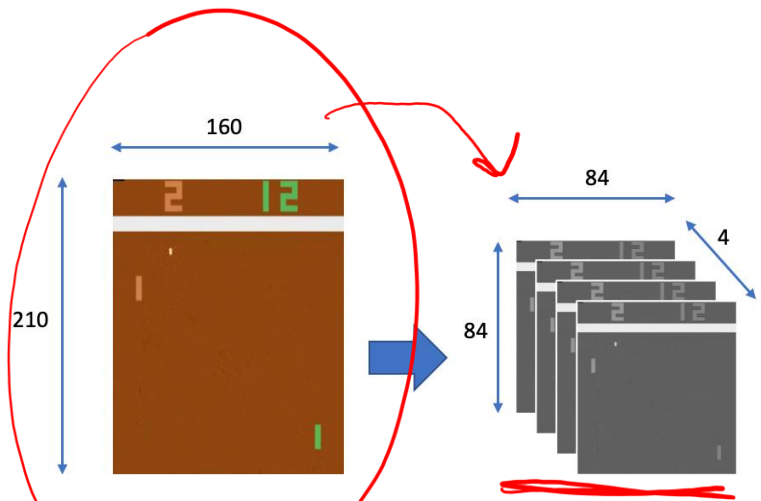
**Markov assumption:**

$r_{t,h}$  and  $s_{t,h+1}$  are conditionally independent of  $(s_{t,1}, a_{t,1}, \dots, s_{t,h-1}, a_{t,h-1})$  given  $s_{t,h}$


Goal: maximize  $\sum_{t=1}^T \sum_{h=1}^{\tau_t} R(s_{t,h}, a_{t,h})$

$\tau_t$  ← length of episode


# From Observations to States



Stacking recent observations



Recurrent neural network




Hidden Markov model

# Regret (Episodic Setting)

Policy : mapping from state to action  
(action distribution)

$$\text{Regret} = \underbrace{\max_{\pi^*} \mathbb{E}^{\pi^*} \left[ \sum_{t=1}^T \sum_{h=1}^{\tilde{\tau}_t} R(\tilde{s}_{t,h}, \pi^*(\tilde{s}_{t,h})) \right]}_{\text{Benchmark}} - \sum_{t=1}^T \sum_{h=1}^{\tau_t} R(s_{t,h}, a_{t,h})$$




















# Example: Racing


- A robot car wants to travel far, quickly
- Three states: **Cool**, **Warm**, Overheated
- Two actions: *Slow*, *Fast*
- Going faster gets double reward

$$\left\{ \begin{array}{l} R(s,a) \\ P(s'|s,a) \end{array} \right\} \leftarrow$$



# Example: Racing

$s$	$a$	$s'$	$P(s' s, a)$	$R(s, a)$
	Slow		1.0	+1
	Fast		0.5	+2
	Fast		0.5	+2
	Slow		0.5	+1
	Slow		0.5	+1
	Fast		1.0	-10
	(end)		1.0	0



# Formulations

- Interaction Protocol
  - Fixed-Horizon
  - Variable-Horizon
- Performance Metric
  - Total Reward
  - Discounted Reward
- Policy
  - Markov policy
  - Stationary policy

Horizon = Length of an episode

# Interaction Protocols (1/2): Fixed-Horizon

Horizon length is a fixed number  $H$

$h \leftarrow 1$

Observe initial state  $s_1 \sim \rho$

**While**  $h \leq H$ :

Choose action  $a_h$

Observe reward  $r_h$  with  $\mathbb{E}[r_h] = R(s_h, a_h)$

Observe next state  $s_{h+1} \sim P(\cdot | s_h, a_h)$

**Examples:** games with a fixed number of time

# Interaction Protocols (2/2): Variable-Horizon

The learner interacts with the environment until reaching **terminal states**  $\mathcal{T} \subset \mathcal{S}$

$h \leftarrow 1$

Observe initial state  $s_1 \sim \rho$

**While**  $s_h \notin \mathcal{T}$ :

    Choose action  $a_h$

    Observe reward  $r_h$  with  $\mathbb{E}[r_h] = R(s_h, a_h)$

    Observe next state  $s_{h+1} \sim P(\cdot | s_h, a_h)$

$h \leftarrow h + 1$

**Examples:** video games, robotics tasks, personalized recommendations, etc.

# Formulations

- Interaction Protocol
  - Fixed-Horizon
  - Variable-Horizon
- Performance Metric
  - Total Reward
  - Discounted Reward
- Policy
  - Markov policy
  - Stationary policy

Horizon = Length of an episode

# Performance Metric

$\tau$ : the step where the episode ends

**Total Reward:**

$$\sum_{h=1}^{\tau} r_h$$

**Discounted Total Reward:**

$$\sum_{h=1}^{\tau} \gamma^{h-1} r_h$$

$$\begin{aligned} r_h &\in [-1, 1] \\ \Rightarrow \left| \text{Discounted total reward} \right| &\leq 1 + \gamma + \gamma^2 + \dots + \gamma^{\tau-1} \\ &\leq \frac{1}{1-\gamma} \quad \checkmark \end{aligned}$$

$$\begin{aligned} \gamma &\approx 0.9 \rightarrow \frac{1}{1-\gamma} = 10 \quad \text{equivalent} \\ &\approx 0.99 \rightarrow \frac{1}{1-\gamma} = 100 \\ \gamma &\in [0, 1): \text{ discount factor} \end{aligned}$$

$$r_1 + \gamma r_2 + \gamma^2 r_3 + \dots +$$

Due to discounting, the future reward starting from any state is always upper bounded by  $\frac{\text{range of } r}{1-\gamma}$ , even if the episode length is very very long.

Without discounting, the range of future reward could be unbounded  $\rightarrow$  making it hard to optimize

There is a potential mismatch between our ultimate goal and what we really optimized.

# Formulations

- Interaction Protocol
  - Fixed-Horizon
  - Variable-Horizon
- Performance Metric
  - Total Reward
  - Discounted Reward
- Policy
  - Markov policy
  - Stationary policy



# Policy for MDPs

$$a_h \sim \pi(\cdot | s_1, a_1, s_2, a_2, \dots, s_h)$$

history-dependent

## Markov Policy


$$a_h \sim \pi_h(\cdot | s_h)$$
$$a_h = \pi_h(s_h)$$

For **fixed-horizon** setting, there exists an optimal policy in this class

## Stationary Policy

$$a_h \sim \pi(\cdot | s_h)$$
$$a_h = \pi(s_h)$$

For **variable-horizon** settings, there exists an optimal policy in this class



slow: +1  
fast: 2

Markov Policy = Stationary Policy where the state is augmented with **the timestep**.

A **stationary policy** specifies

$$\pi(\text{Slow} \mid \text{Cool})$$

$$\pi(\text{Fast} \mid \text{Cool})$$

$$\pi(\text{Slow} \mid \text{Warm})$$

$$\pi(\text{Fast} \mid \text{Warm})$$

A **Markov policy** specifies


$$\pi_h(\text{Slow} \mid \text{Cool})$$

$$\pi_h(\text{Fast} \mid \text{Cool})$$

$$\pi_h(\text{Slow} \mid \text{Warm})$$

$$\pi_h(\text{Fast} \mid \text{Warm})$$

$$\forall h$$



# **Value Iteration**

(Fixed-Horizon + Total-Reward)


# Two Tasks

**Policy Evaluation:** Calculate the expected total reward of a given policy

What is the expected total reward for the policy  $\pi(\text{cool}) = \text{fast}$ ,  $\pi(\text{warm}) = \text{slow}$ ?

**Policy Optimization:** Find the best policy


What is the policy that achieves the highest expected total reward?



# Value Iteration for Policy Evaluation

$$V_h^\pi(s) = \mathbb{E}^\pi \left[ \sum_{k=h}^H R(s_k, a_k) \mid s_h = s \right]$$

Fix  $\pi$   $\pi_h(a|s)$



State transition:  $P(s'|s, a)$

Reward:  $R(s, a)$

$$Q_h^\pi(s, a) = \mathbb{E}^\pi \left[ \sum_{k=h}^H R(s_k, a_k) \mid (s_h, a_h) = (s, a) \right]$$

$$V_h^\pi(s) = \mathbb{E}^\pi \left[ \sum_{k=h}^H R(s_k, a_k) \mid s_h = s \right]$$

**Backward induction:**

$$V_{H+1}^\pi(s) = 0 \quad \forall s$$

For  $h = H, \dots, 1$ : for all  $s, a$

$$Q_h^\pi(s, a) = R(s, a) + \sum_{s'} P(s'|s, a) V_{h+1}^\pi(s')$$

Expected total reward of  $\pi$  from step  $h + 1$

$$V_h^\pi(s) = \sum_a \pi_h(a|s) Q_h^\pi(s, a)$$

# Bellman Equation

$Q_h^\pi$  is called “the state-action value functions of policy  $\pi$ ”  
 $V_h^\pi$  is called “the state value function of policy  $\pi$ ”  
Both can be just called “**value functions**”

$$Q_h^\pi(s, a) = R(s, a) + \sum_{s'} P(s'|s, a) V_{h+1}^\pi(s')$$

$$V_h^\pi(s) = \sum_a \pi_h(a|s) Q_h^\pi(s, a)$$

or

$$\underline{Q_h^\pi(s, a)} = R(s, a) + \sum_{s', a'} P(s'|s, a) \pi_{h+1}(a'|s') \underline{Q_{h+1}^\pi(s', a')}$$

or

$$\underline{V_h^\pi(s)} = \sum_a \pi_h(a|s) \left( R(s, a) + \sum_{s'} P(s'|s, a) \underline{V_{h+1}^\pi(s')} \right)$$

# The Meaning of Bellman Equations

## Definitions

$\mathcal{R}, \mathcal{P}, \pi$

$$\underline{Q_h^\pi(s, a)} \triangleq \mathbb{E}^\pi \left[ \sum_{k=h}^H R(s_k, a_k) \mid (s_h, a_h) = (s, a) \right]$$

$$\underline{V_h^\pi(s)} \triangleq \mathbb{E}^\pi \left[ \sum_{k=h}^H R(s_k, a_k) \mid s_h = s \right]$$

## Relations (Bellman Equations)

$$Q_h^\pi(s, a) = R(s, a) + \sum_{s'} P(s'|s, a) V_{h+1}^\pi(s')$$

$$V_h^\pi(s) = \sum_a \pi_h(a|s) Q_h^\pi(s, a)$$

## Calculation (VI)


Calculate

$Q_h^\pi(s, a), V_h^\pi(s) \forall s, a$   
from  $h = H$  to  $h = 1$

Based on Dynamic Programming

$V_1^*(s)$

# Value Iteration for Policy Optimization



- ✓ State transition:  $P(s'|s, a)$
- ✓ Reward:  $R(s, a)$

$$\underline{Q_h^*(s, a)} = \max_{\pi \in \text{Markov Policy}} \mathbb{E}^\pi \left[ \sum_{k=h}^H R(s_k, a_k) \mid (s_h, a_h) = (s, a) \right]$$

→  $\underline{V_h^*(s)} = \max_{\pi \in \text{Markov Policy}} \mathbb{E}^\pi \left[ \sum_{k=h}^H R(s_k, a_k) \mid s_h = s \right]$

**Backward induction:**

$\pi(a|s)$ : distributions over action

$\pi(s)$ : action

$$\underline{V_{H+1}^*(s)} = 0 \quad \forall s$$

For  $h = H, \dots, 1$ : for all  $s, a$

$$\underline{Q_h^*(s, a)} = R(s, a) + \underbrace{\sum_{s'} P(s'|s, a) V_{h+1}^*(s')}_{\text{Expected optimal total reward from step } h+1}$$















Expected optimal total  
reward from step  $h+1$

$$\underline{V_h^*(s)} = \max_a Q_h^*(s, a)$$

$$\pi_h^*(s) = \operatorname{argmax}_a Q_h^*(s, a)$$



# Exercise

$s$	$a$	$s'$	$P(s' s, a)$	$R(s, a)$
	Slow		1.0	+1
	Fast		0.5	+2
	Fast		0.5	+2
	Slow		0.5	+1
	Slow		0.5	+1
	Fast		1.0	-10
	(end)		1.0	0

Assume  $H = 3$

$$Q_3^*(s, a)$$

$$Q_3^*(\text{cool}, \text{slow}) = 1$$

$$Q_3^*(\text{cool}, \text{fast}) = 2$$

$$Q_3^*(\text{warm}, \text{slow}) = 1$$

$$Q_3^*(\text{warm}, \text{fast}) = -10$$

$$V_3^*(s)$$

$$V_3^*(\text{cool}) = 2$$

$$\pi_3^*(\text{cool}) = \text{fast}$$

$$V_3^*(\text{warm}) = 1$$

$$\pi_3^*(\text{warm}) = \text{slow}$$

$$Q_2^*(s, a) = R(s, a) + \sum_{s'} P(s'|s, a) V_3^*(s')$$

$$Q_2^*(\text{cool}, \text{slow}) = 1 + V_3^*(\text{cool}) = 1 + 2 = 3$$

$$Q_2^*(\text{cool}, \text{fast}) = 2 + \frac{1}{2} V_3^*(\text{cool}) + \frac{1}{2} V_3^*(\text{warm}) = 3.5$$

$$Q_2^*(\text{warm}, \text{slow}) = 1 + \frac{1}{2} V_3^*(\text{cool}) + \frac{1}{2} V_3^*(\text{warm}) = 2.5$$

$$Q_2^*(\text{warm}, \text{fast}) = -10 + V_3^*(\text{overheat}) = -10$$

$$V_2^*(s)$$

$$V_2^*(\text{cool}) = 3.5$$

$$\pi_2^*(\text{cool}) = \text{fast}$$

$$V_2^*(\text{warm}) = 2.5$$

$$\pi_2^*(\text{warm}) = \text{slow}$$

# Bellman Optimality Equation

$Q_h^*$  : optimal state-action value functions

$V_h^*$  : optimal state value functions

or “**optimal value functions**”

$$Q_h^*(s, a) = R(s, a) + \sum_{s'} P(s'|s, a) V_{h+1}^*(s')$$

✓  $V_h^*(s) = \max_a Q_h^*(s, a)$

or

$$Q_h^*(s, a) = R(s, a) + \sum_{s'} P(s'|s, a) \left( \max_{a'} Q_{h+1}^*(s', a') \right)$$

or


$$V_h^*(s) = \max_a \left( R(s, a) + \sum_{s'} P(s'|s, a) V_{h+1}^*(s') \right)$$

$$\pi_h^*(s) = \operatorname{argmax}_a Q_h^*(s, a)$$

# **Value Iteration**

(Variable-Horizon + Discounted Reward)

# Value Iteration for Policy Evaluation



State transition:  $P(s'|s, a)$

Reward:  $R(s, a)$

$$Q_i^\pi(s, a) = R(s, a) + \gamma R(s_2, a_2) + \gamma^2 R(s_3, a_3) + \dots + \gamma^{i-1} R(s_i, a_i)$$

$$Q_i^\pi(s, a) = \mathbb{E}^\pi \left[ \sum_{h=1}^i \gamma^{h-1} R(s_h, a_h) \mid (s_1, a_1) = (s, a) \right]$$

$$V_i^\pi(s) = \mathbb{E}^\pi \left[ \sum_{h=1}^i \gamma^{h-1} R(s_h, a_h) \mid s_1 = s \right]$$

Handwritten expansion of the expectation:

$$R(s, a) + \gamma (R(s_2, a_2) + \gamma R(s_3, a_3) + \dots + \gamma^{i-1} R(s_i, a_i))$$

*i-1 steps remaining*

$$Q^\pi(s, a) = Q_\infty^\pi(s, a) \quad V^\pi(s) = V_\infty^\pi(s)$$

$$V_0^\pi(s) = 0 \quad \forall s$$

For  $i = 1, 2, 3, \dots$  for all  $s, a$

$$Q_i^\pi(s, a) = R(s, a) + \gamma \sum_{s'} P(s'|s, a) V_{i-1}^\pi(s')$$

$$V_i^\pi(s) = \sum_a \pi(a|s) Q_i^\pi(s, a)$$

If  $|Q_i^\pi(s, a) - Q_{i-1}^\pi(s, a)| \leq \epsilon$  for all  $s, a$ : **terminate**

# Bellman Equation

$$Q^{\pi}(s, a) = R(s, a) + \gamma \sum_{s'} P(s'|s, a) V^{\pi}(s')$$

$$V^{\pi}(s) = \sum_a \pi(a|s) Q^{\pi}(s, a)$$

or

$$Q^{\pi}(s, a) = R(s, a) + \gamma \sum_{s', a'} P(s'|s, a) \pi(a'|s') Q^{\pi}(s', a')$$

or

$$V^{\pi}(s) = \sum_a \pi(a|s) \left( R(s, a) + \gamma \sum_{s'} P(s'|s, a) V^{\pi}(s') \right)$$

# The Meaning of Bellman Equations

## Definitions

$$Q^{\pi}(s, a) = \mathbb{E}^{\pi} \left[ \sum_{h=1}^{\infty} \gamma^{h-1} R(s_h, a_h) \mid (s_1, a_1) = (s, a) \right]$$

$$V^{\pi}(s) = \mathbb{E}^{\pi} \left[ \sum_{h=1}^{\infty} \gamma^{h-1} R(s_h, a_h) \mid s_1 = s \right]$$

## Relations (Bellman Equations)

$$Q^{\pi}(s, a) = R(s, a) + \gamma \sum_{s'} P(s'|s, a) V^{\pi}(s')$$

$$V^{\pi}(s) = \sum_a \pi(a|s) Q^{\pi}(s, a)$$

## Calculation (VI)

Calculate  
 $Q_i^{\pi}(s, a), V_i^{\pi}(s) \forall s, a$   
for  $i = 1, 2, \dots$   
until terminated

# The Quality of $Q_i^\pi(s, a)$ when VI Terminates

Unanswered questions:

1. Will VI (for policy evaluation) always terminate?
2. At termination, we know  $\max_{s,a} |Q_i^\pi(s, a) - Q_{i-1}^\pi(s, a)| \leq \epsilon$ ,  
but our goal is to approximate  $Q^\pi(s, a)$ .

What can we say about  $\max_{s,a} |Q_i^\pi(s, a) - Q^\pi(s, a)|$ ?

# The Quality of $Q_i^\pi(s, a)$ when VI Terminates

Let  $f: \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$  be **any** function. Define

$$\text{BellmanError}(f) = \max_{s,a} \left| f(s, a) - \left( R(s, a) + \gamma \sum_{s',a'} P(s'|s, a) \pi(a'|s') f(s', a') \right) \right|$$

$$\text{ValueError}(f) = \max_{s,a} |f(s, a) - Q^\pi(s, a)|$$

## Theorem


$$\text{ValueError}(f) \leq \frac{\text{BellmanError}(f)}{1 - \gamma}$$

With this theorem, we can argue the quality of  $Q_i^\pi(s, a)$  when VI terminates through the following:

1. Prove that when VI terminates,  $\text{BellmanError}(Q_i^\pi) \leq \epsilon$
2. Using the theorem, we get  $\text{ValueError}(Q_i^\pi) \leq \frac{\epsilon}{1-\gamma}$



# Value Iteration for Policy Optimization



State transition:  $P(s'|s, a)$

Reward:  $R(s, a)$

$$Q_i^*(s, a) = \max_{\pi} \mathbb{E}^{\pi} \left[ \sum_{h=1}^i \gamma^{h-1} R(s_h, a_h) \mid (s_0, a_0) = (s, a) \right]$$

$$V_i^*(s) = \max_{\pi} \mathbb{E}^{\pi} \left[ \sum_{h=1}^i \gamma^{h-1} R(s_h, a_h) \mid s_0 = s \right]$$

$$Q^*(s, a) = Q_{\infty}^*(s, a) \quad V^*(s) = V_{\infty}^*(s)$$

$$V_0^*(s) = 0 \quad \forall s$$

For  $i = 1, 2, 3, \dots$  for all  $s, a$

$$Q_i^*(s, a) = R(s, a) + \gamma \sum_{s'} P(s'|s, a) V_{i-1}^*(s')$$

$$V_i^*(s) = \max_a Q_i^*(s, a)$$

If  $|Q_i^*(s, a) - Q_{i-1}^*(s, a)| \leq \epsilon$  for all  $s, a$  : **terminate**

# Bellman Optimality Equation

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$

$$Q^*(s, a) = R(s, a) + \gamma \sum_{s'} P(s'|s, a) V^*(s')$$

$$V^*(s) = \max_a Q^*(s, a)$$

or

$$Q^*(s, a) = R(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q^*(s', a')$$

or

$$V^*(s) = \max_a \left( R(s, a) + \gamma \sum_{s'} P(s'|s, a) V^*(s') \right)$$

# The Solution Quality when VI Terminates

Unanswered questions:

1. Will VI (for policy optimization) always terminate?

2. At termination, we know  $\max_{s,a} |Q_i^*(s, a) - Q_{i-1}^*(s, a)| \leq \epsilon$ ,

What can we say about  $\max_{s,a} |Q_i^*(s, a) - Q^*(s, a)|$ ?

3. And what can we say about the **performance of the greedy policy  $\hat{\pi}$**

defined as  $\hat{\pi}(a|s) = \mathbb{I} \left[ a = \operatorname{argmax}_{a'} Q_i^*(s, a') \right]$ ? or simply  $\hat{\pi}(s) = \operatorname{argmax}_{a'} Q_i^*(s, a')$

# The Solution Quality when VI Terminates (1/2)

Let  $f: \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$  be **any** function. Define

$$\text{BellmanError}(f) = \max_{s,a} \left| f(s,a) - \left( R(s,a) + \gamma \sum_{s'} P(s'|s,a) \max_{a'} f(s',a') \right) \right|$$

$$\text{ValueError}(f) = \max_{s,a} |f(s,a) - Q^*(s,a)|$$

## Theorem

$$\text{ValueError}(f) \leq \frac{\text{BellmanError}(f)}{1 - \gamma}$$

# The Solution Quality when VI Terminates (2/2)

Let  $f: \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$  be **any** function. Define

$$\pi_f(s) = \operatorname{argmax}_a f(s, a)$$

## Theorem

$$V^*(\rho) - V^{\pi_f}(\rho) \leq \frac{2}{1-\gamma} \text{ValueError}(f)$$

Combining the two theorems, we know that when VI (for policy optimization) terminates,

$$V^*(\rho) - V^{\hat{\pi}}(\rho) \leq \frac{2}{1-\gamma} \text{ValueError}(Q_i^*) \leq \frac{2}{(1-\gamma)^2} \text{BellmanError}(Q_i^*) \leq \frac{2\epsilon}{(1-\gamma)^2}$$

where  $\hat{\pi}(s) = \operatorname{argmax}_a Q_i^*(s, a)$

# **Policy Iteration**

# Policy Iteration

## Policy Iteration

For  $i = 1, 2, \dots$

$$\forall s, \quad \pi_i(s) \leftarrow \operatorname{argmax}_a Q^{\pi_{i-1}}(s, a)$$

**Theorem (monotonic improvement).** Policy Iteration ensures

$$\forall s, a, \quad Q^{\pi_i}(s, a) \geq Q^{\pi_{i-1}}(s, a)$$

When converged (i.e.,  $\pi_i = \pi_{i-1}$ ), we have  $\pi_i = \pi^*$ .

(We will prove this later.)

# Generalized Policy Iteration

$N = \infty \Rightarrow$  Policy Iteration

$N = 1 \Rightarrow$  Value Iteration for policy optimization

For  $i = 1, 2, \dots$

$$\pi_i(s) = \max_a Q_i(s, a) \quad \leftarrow \text{Policy update}$$

$$Q \leftarrow Q_i$$

Repeat for  $N$  times:

$$Q(s, a) \leftarrow R(s, a) + \gamma \sum_{s', a'} P(s' | s, a) \pi_i(a' | s') Q(s', a')$$

$$Q_{i+1} \leftarrow Q$$

Value update

**Notice:** in value iteration for PO, there may not exist a policy  $\pi$  such that  $Q_i = Q^\pi$

In contrast, in policy iteration we have  $Q_i = Q^{\pi_{i-1}}$

VI for PO can be viewed as PI **with incomplete policy evaluation**



# Summary

- Value Iteration for Policy Optimization (VI for PO)
  - Is essentially a **dynamic programming** algorithm
  - Finds the value functions of the optimal policy
- Value Iteration for Policy Evaluation (VI for PE)
  - Also a **dynamic programming** algorithm
  - Finds the value functions of the given policy
- Policy Iteration (PI)
  - An iterative policy improvement algorithm
  - Each iteration involves a policy evaluation subtask
- VI for PO and PI can be viewed as special cases of Generalized PI

# **Performance Difference Lemma**

# Recall: Regret

$$\text{Regret} = \max_{\pi^*} \mathbb{E}^{\pi^*} \left[ \sum_{t=1}^T \sum_{h=1}^{\tilde{\tau}_t} R(\tilde{s}_{t,h}, \pi^*(\tilde{s}_{t,h})) \right] - \sum_{t=1}^T \sum_{h=1}^{\tau_t} R(s_{t,h}, a_{t,h})$$

$$\mathbb{E}[\text{Regret}] = \mathbb{E} \left[ \sum_{t=1}^T (V_1^*(s_{t,1}) - V_1^{\pi_t}(s_{t,1})) \right]$$

$$= \mathbb{E} \left[ \sum_{t=1}^T (V_1^*(\rho) - V_1^{\pi_t}(\rho)) \right] \quad V_1^{\pi}(\rho) \triangleq \mathbb{E}_{s \sim \rho} [V_1^{\pi}(s)]$$

# Unanswered Questions

- For an estimation  $\hat{Q}(s, a) \approx Q^*(s, a)$  with error, how can we bound

$$V^*(\rho) - V^{\hat{\pi}}(\rho) \quad \text{where } \hat{\pi}(s) = \operatorname{argmax}_a \hat{Q}(s, a)?$$

- How to show that Policy Iteration leads to monotonic policy improvement?
- Also, how are these methods related to the third challenge of online RL: credit assignment?

# Performance Difference Lemma

For any two stationary policies  $\pi'$  and  $\pi$  in the discounted setting,

$$\begin{aligned}\mathbb{E}_{s \sim \rho} [V^{\pi'}(s)] - \mathbb{E}_{s \sim \rho} [V^{\pi}(s)] &= \sum_{s,a} d_{\rho}^{\pi'}(s) (\pi'(a|s) - \pi(a|s)) Q^{\pi}(s, a) \\ &= \sum_{s,a} d_{\rho}^{\pi'}(s, a) (Q^{\pi}(s, a) - V^{\pi}(s))\end{aligned}$$

$$d_{\rho}^{\pi}(s) \triangleq \mathbb{E}^{\pi} \left[ \sum_{h=1}^{\infty} \gamma^{h-1} \mathbb{I}\{s_h = s\} \mid s_1 \sim \rho \right] \quad \text{Discounted occupancy measure on state } s$$

$$d_{\rho}^{\pi}(s, a) \triangleq \mathbb{E}^{\pi} \left[ \sum_{h=1}^{\infty} \gamma^{h-1} \mathbb{I}\{(s_h, a_h) = (s, a)\} \mid s_1 \sim \rho \right]$$

# Performance Difference Lemma

We can also swap the roles of  $\pi'$  and  $\pi$  and apply the same lemma

$$\mathbb{E}_{s \sim \rho}[V^\pi(s)] - \mathbb{E}_{s \sim \rho}[V^{\pi'}(s)] = \sum_{s,a} d_\rho^\pi(s) (\pi(a|s) - \pi'(a|s)) Q^{\pi'}(s, a)$$

$$\stackrel{\times (-1)}{\Rightarrow} \mathbb{E}_{s \sim \rho}[V^{\pi'}(s)] - \mathbb{E}_{s \sim \rho}[V^\pi(s)] = \sum_{s,a} d_\rho^\pi(s) (\pi'(a|s) - \pi(a|s)) Q^{\pi'}(s, a)$$

||

Original version:

$$\mathbb{E}_{s \sim \rho}[V^{\pi'}(s)] - \mathbb{E}_{s \sim \rho}[V^\pi(s)] = \sum_{s,a} d_\rho^{\pi'}(s) (\pi'(a|s) - \pi(a|s)) Q^\pi(s, a)$$

# Performance Difference Lemma (Fixed-Horizon)

For any two Markov policies  $\pi'$  and  $\pi$  in the fixed-horizon setting,

$$\begin{aligned}\mathbb{E}_{s_1 \sim \rho} [V_1^{\pi'}(s_1)] - \mathbb{E}_{s_1 \sim \rho} [V_1^{\pi}(s_1)] &= \sum_{h=1}^H \sum_{s,a} d_{\rho,h}^{\pi'}(s) (\pi'_h(a|s) - \pi_h(a|s)) Q_h^{\pi}(s,a) \\ &= \sum_{h=1}^H \sum_{s,a} d_{\rho,h}^{\pi'}(s,a) (Q_h^{\pi}(s,a) - V_h^{\pi}(s))\end{aligned}$$

$$d_{\rho,h}^{\pi}(s) \triangleq \mathbb{E}^{\pi}[\mathbb{I}\{s_h = s\} \mid s_1 \sim \rho] = \mathbb{P}^{\pi}(s_h = s \mid s_1 \sim \rho)$$

$$d_{\rho,h}^{\pi}(s,a) \triangleq \mathbb{E}^{\pi}[\mathbb{I}\{(s_h, a_h) = (s,a)\} \mid s_1 \sim \rho] = \mathbb{P}^{\pi}((s_h, a_h) = (s,a) \mid s_1 \sim \rho)$$

# The Meaning of Performance Difference Lemma

It tells us how **credit** are **assigned** to each state/step

The sub-optimality of a policy  $\pi$ :

$$\mathbb{E}_{s \sim \rho}[V^*(s)] - \mathbb{E}_{s \sim \rho}[V^\pi(s)]$$

If  $\pi$  is highly sub-optimal, then we can always find

- 1) An  $(s, a)$ -pair on the path of  $\pi$  such that  $V^*(s) - Q^*(s, a)$  is positive and large
- 2) An  $(s, a)$ -pair on the path of  $\pi^*$  such that  $Q^\pi(s, a) - V^\pi(s)$  is positive and large

$$= \sum_{s,a} d_\rho^\pi(s) (\pi^*(a|s) - \pi(a|s)) Q^{\pi^*}(s, a)$$

$$= \sum_{s,a} d_\rho^\pi(s, a) (V^*(s) - Q^*(s, a)) \quad \checkmark$$


$$= \sum_{s,a} d_\rho^{\pi^*}(s) (\pi^*(a|s) - \pi(a|s)) Q^\pi(s, a)$$

$$= \sum_{s,a} d_\rho^{\pi^*}(s, a) (Q^\pi(s, a) - V^\pi(s)) \quad \checkmark$$



A game tree for the 'X' player, where the 'O' player always plays in the **first** available cell (from left to right, top to bottom).

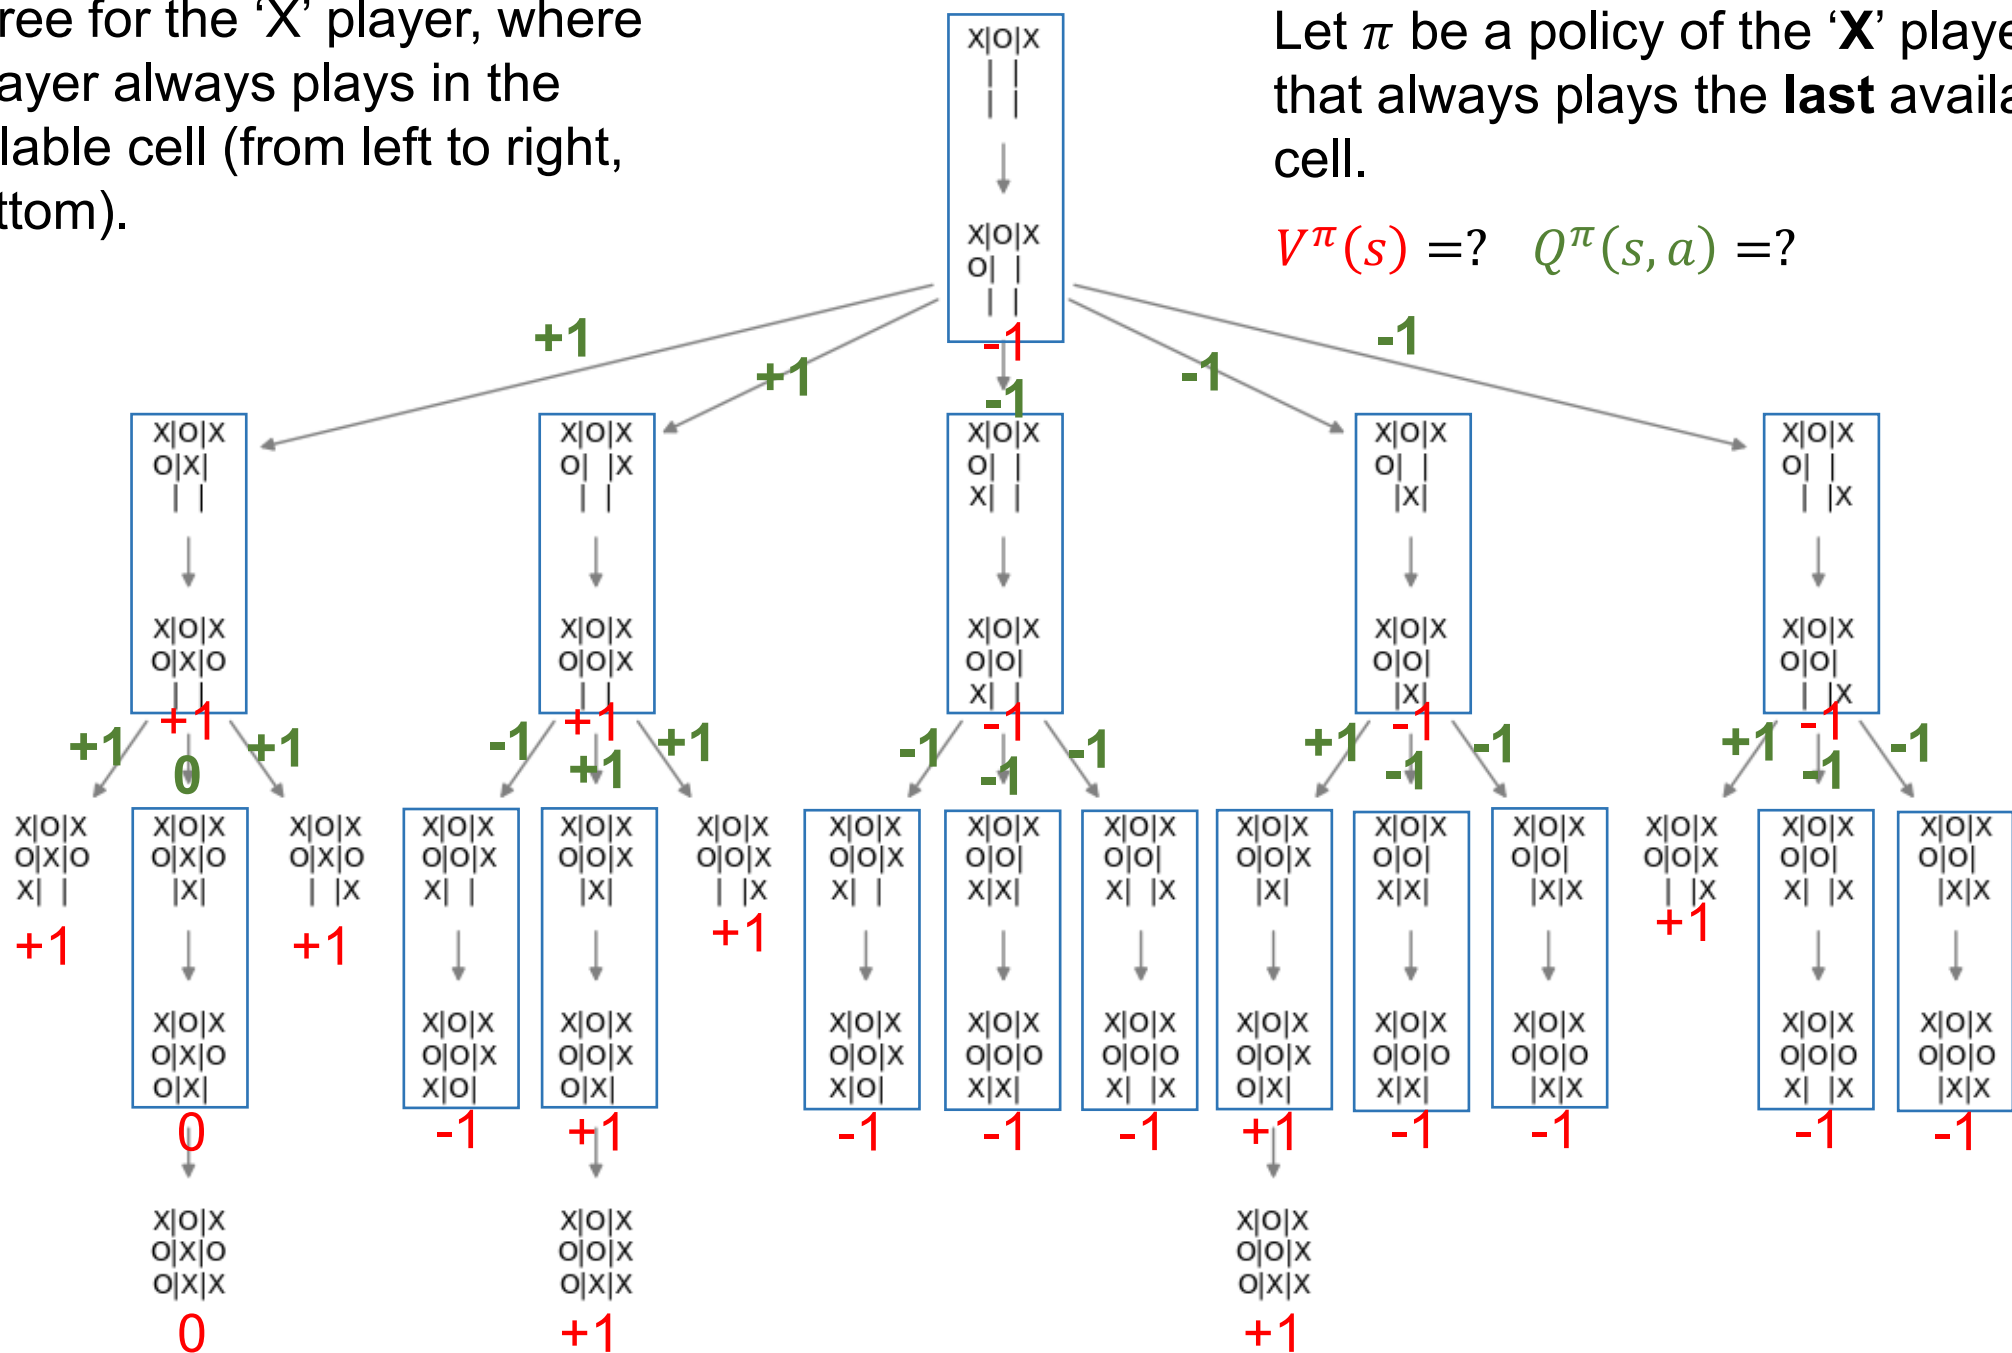
$V^*(s) = ?$     $Q^*(s, a) = ?$



A game tree for the 'X' player, where the 'O' player always plays in the **first** available cell (from left to right, top to bottom).

Let  $\pi$  be a policy of the 'X' player that always plays the **last** available cell.

$V^\pi(s) = ? \quad Q^\pi(s, a) = ?$



# **Proof (Sketch) of Performance Difference Lemma**

h



# Unanswered Question 1

**Suboptimality  $\leq (1 - \gamma)^{-1}$  Value Error**

Let  $f: \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$  be **any** function

If

$$|f(s, a) - Q^*(s, a)| \leq \epsilon \quad \forall s, a$$

then

$$\underbrace{V^*(s) - V^{\pi_f}(s)} \leq \frac{2\epsilon}{1 - \gamma} \quad \forall s$$

where  $\pi_f(s) = \operatorname{argmax}_a f(s, a)$



## Unanswered Question 2

Policy Iteration ensures

$$\forall s, a, \quad Q^{\pi_i}(s, a) \geq Q^{\pi_{i-1}}(s, a)$$

When converged (i.e.,  $\pi_i = \pi_{i-1}$ ), we have  $\pi_i = \pi^*$ .



$$\pi_i = \pi_{i-1}$$

$$\Rightarrow \pi_i(s) = \operatorname{argmax}_a Q^{\pi_i}(s, a)$$

$$\Rightarrow Q^{\pi_i}(s, a) = R(s, a) + \gamma \sum_{s', a'} P(s'|s, a) \pi_i(a'|s') Q^{\pi_i}(s', a') = R(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q^{\pi_i}(s', a')$$

$$\Rightarrow Q^{\pi_i} \text{ satisfies the Bellman optimality equation}$$

$$\Rightarrow \text{BellmanError}(Q^{\pi_i}) = 0$$

$$\Rightarrow Q^{\pi_i}(s, a) = Q^*(s, a) \text{ by the “ValueError} \leq \frac{1}{1-\gamma} \text{BellmanError” lemma on Page 38}$$

$$\Rightarrow \pi_i(s) = \operatorname{argmax}_a Q^*(s, a) = \pi^*(s).$$

# Recap: MDP

- Definitions of  $Q^\pi(s, a)$ ,  $V^\pi(s)$ ,  $Q^*(s, a)$ ,  $V^*(s)$
- Bellman equations (related to dynamic programming)
- Value Iteration to approximate  $Q^\pi(s, a)/V^\pi(s)$  or  $Q^*(s, a)/V^*(s)$
- Policy Iteration to find  $\pi^*$  --- involving  $Q^\pi(s, a)/V^\pi(s)$  approximation
- Unified by Generalized Policy Iteration
- Performance difference lemma to decompose  $\mathbb{E}_{s \sim \rho} [V^{\pi'}(s)] - \mathbb{E}_{s \sim \rho} [V^\pi(s)]$ 
  - Credit assignment
  - $= \sum_{s,a} d_\rho^\pi(s, a) (V^{\pi'}(s) - Q^{\pi'}(s, a))$  (helpful in analyzing VI by letting  $\pi' = \pi^*$ )
  - $= \sum_{s,a} d_\rho^{\pi'}(s, a) (Q^\pi(s, a) - V^\pi(s))$  (helpful in analyzing PI by letting  $\pi' = \pi_{i+1}$ )

# Next

- Our discussion indicates there are two potential ways to find optimal policy
  - Value-Iteration-based: approximate  $\hat{Q}(s, a) \approx Q^*(s, a)$  and let  $\hat{\pi}(s) = \operatorname{argmax}_a \hat{Q}(s, a)$
  - Policy-Iteration-based: approximate  $\hat{Q}(s, a) \approx Q^\pi(s, a)$  and let  $\hat{\pi}(s) = \operatorname{argmax}_a \hat{Q}(s, a)$
  - ... or something in between (based on generalized policy iteration)
- RL algorithms we will discuss:
  - Performing approximate VI or PI using samples