

Approximate Policy Iteration and Variants

Chen-Yu Wei

Policy Iteration

For $k = 1, 2, \dots$

Calculate $Q^{\pi_k}(s, a) \quad \forall s, a$

$\pi_{k+1}(s) = \operatorname{argmax}_a Q^{\pi_k}(s, a) \quad \forall s$

Asynchronous Policy Iteration

For $k = 1, 2, \dots$

Pick any state \hat{s}

Calculate $Q^{\pi_k}(\hat{s}, a) \quad \forall a$

$\pi_{k+1}(\hat{s}) = \underset{a}{\operatorname{argmax}} Q^{\pi_k}(\hat{s}, a)$

and $\pi_{k+1}(s) = \pi_k(s) \quad \forall s \neq \hat{s}$

$$\underbrace{Q^{\pi_{k+1}}(s, a)}_{\geq Q^{\pi_k}(s, a)} \quad \forall s, a$$

$$\mathbb{E}_{s \sim p} [V^{\pi_{k+1}}(s)] - \mathbb{E}_{s \sim p} [V^{\pi_k}(s)]$$

$$= \sum_{s, a} d^{\pi_{k+1}}(s) \left(\underbrace{\pi_{k+1}(a|s) - \pi_k(a|s)}_{\geq 0} \right) Q^{\pi_k}(s, a)$$

$$= \sum_a d^{\pi_{k+1}}(\hat{s}) \left(\underbrace{\pi_{k+1}(a|\hat{s}) - \pi_k(a|\hat{s})}_{\geq 0} \right) Q^{\pi_k}(\hat{s}, a)$$

$$\geq 0$$

[If we want this to be positive & large]

Asynchronous Policy Iteration

- To improve policy, we may just evaluate Q^{π_k} on a particular state s .
- Of course, a **real improvement** is made only when $\exists a \text{ s.t. } \underline{Q^{\pi_k}(s, a) - V^{\pi_k}(s)}$ is large.
- This is **different from Value Iteration**, where ideally, we would like to find Q_{k+1} such that $Q_{k+1}(s, a) \approx R(s, a) + \gamma \mathbb{E}_{s' \sim P(\cdot|s, a)} \left[\max_{a'} Q_k(s', a') \right] \quad \forall s, a$
- VI-based algorithm like DQN usually requires **stronger function approximation** that can generalize to unseen state.

Policy Iteration with Samples

For $k = 1, 2, \dots$

For $i = 1, 2, \dots, N$:

Choose action $a_i \sim \pi_{\theta_k}(\cdot | s_i)$

Receive reward $r_i \sim R(s_i, a_i)$ and $s'_i \sim P(\cdot | s_i, a_i)$

$s_{i+1} = s'_i$ if episode continues, $s_{i+1} \sim \rho$ if episode ends

~~✓ Evaluate $Z_k(s, a) \approx Q^{\pi_{\theta_k}}(s, a)$ for $s = s_1, \dots, s_N$ and all a~~
or $Z_k(s, a) \approx Q^{\pi_{\theta_k}}(s, a) - b_k(s)$ for $s = s_1, \dots, s_N$ and all a

✓ Update θ_{k+1} from θ_k using the estimators $\{Z_k(s_i, a)\}_{i=1}^N$
Using any technique we introduced for policy-based contextual bandits

$$z_k(s, a) = \hat{R}(s, a)$$

$$z_k(s, a) = \frac{r_t + b(s)}{z_k(a|s)} \mathbb{I}(a_t = a)$$
$$\mathbb{E}[z_k(s, a)] \approx \hat{R}(s, a)$$

Data collection

Policy Evaluation

Policy Improvement

Why can we independently optimize the policy on each state?

Essentially treating **states** as **contexts**, but replacing $R(x, a)$ by $\underline{Q^{\pi_{\theta_k}}(s, a)}$

Policy Evaluation

Policy Evaluation

Given: a policy π

Evaluate $V^\pi(s)$ or $Q^\pi(s, a)$

- {
- On-policy policy evaluation:** the learner can execute π to evaluate π
 - Off-policy/offline policy evaluation:** the learner can only execute some $\pi_b \neq \pi$, or can only access some existing dataset to evaluate π

Use cases:

- Approximate policy iteration: $\pi_k(s) = \operatorname{argmax}_a Q^{\pi_{k-1}}(s, a)$
- Estimate the value of a policy before deploying it in the real world, e.g., COVID-related border measures, economic recovery policies, or policy changes in recommendation systems.

Value Iteration for V^π / Q^π

Assuming P, R are known

Input: π

For $k = 1, 2, \dots$

$$\forall s, V_k(s) \leftarrow \sum_a \pi(a|s) \left(R(s, a) + \gamma \sum_{s'} P(s'|s, a) V_{k-1}(s') \right)$$

Input: π

For $k = 1, 2, \dots$

$Q_k \rightarrow Q^\pi$

$$\forall s, a, Q_k(s, a) \leftarrow R(s, a) + \gamma \sum_{s', a'} P(s'|s, a) \pi(a'|s') Q_{k-1}(s', a')$$

On-Policy Policy Evaluation

Value Iteration for PE (given π)

Temporal Difference (TD) Learning for V^π

correct π : collect

max_a

For $k = 1, 2, \dots$

(Q π)

Collect $\{(s_i, a_i, r_i, s'_i)\}_{i=1}^N$ using policy π

$s'_i \sim P(\cdot | s_i, a_i)$

$$\theta_k \leftarrow \theta_{k-1} - \alpha \nabla_\theta \frac{1}{N} \sum_{i=1}^N \left(V_\theta(s_i) - r_i - \gamma V_{\theta_{k-1}}(s'_i) \right)^2 \Bigg|_{\theta=\theta_{k-1}}$$

No target network needed because this is an **on-policy** problem.

Know P, R

This algorithm is also called TD(0)

with sample:

$$V_k(s) \leftarrow \sum_a \pi(a|s) (R(s,a) + \gamma \sum_{s'} P(s'|s,a) V_{k-1}(s'))$$

$$= \sum_a \pi(a|s) R(s,a)$$

$$+ \gamma \sum_a \pi(a|s) \sum_{s'} P(s'|s,a) V_{k-1}(s')$$

$$\min_\theta \sum_{i=1}^N \left(V_\theta(s_i) - (r_i + \gamma V_{\theta_{k-1}}(s'_i)) \right)^2$$

Temporal Difference (TD) Learning for Q^π

For $k = 1, 2, \dots$

Collect $\{(s_i, a_i, r_i, s'_i)\}_{i=1}^N$ using policy π

$$\theta_k \leftarrow \theta_{k-1} - \alpha \nabla_{\theta} \frac{1}{N} \sum_{i=1}^N \left(Q_{\theta}(s_i, a_i) - r_i - \gamma \sum_a \pi(a|s'_i) Q_{\theta_{k-1}}(s'_i, a') \right)^2 \Bigg|_{\theta=\theta_{k-1}}$$

No target network needed because this is an on-policy problem.

Monte Carlo Estimation

Start from $(s_1, a_1) = (\hat{s}, \hat{a})$ and execute policy π until the episode ends and obtain trajectory

$$s_1 = \hat{s}, a_1 = \hat{a}, r_1, s_2, a_2, r_2, \dots, s_\tau, a_\tau, r_\tau$$

Let $G = \sum_{h=1}^{\tau} \gamma^{h-1} r_h$

$$E[G] = Q^\pi(\hat{s}, \hat{a})$$

G is an unbiased estimator for $Q^\pi(\hat{s}, \hat{a})$

MC estimator: unbiased, higher variance

TD estimator: biased, lower variance

$$\hat{Q}(s, a)$$

$$Q^\pi(s, a) = E \left[\left(\sum_{h=1}^{\infty} \gamma^{h-1} r_h \mid (s_1, a_1) = (s, a) \right) \right]$$

bandit

$$r$$

$$\hat{R}(x_t, a_t)$$

$$z(\hat{s}, \hat{a}) = \frac{G \mathbb{I}\{a=\hat{a}\}}{\pi(\hat{a}|\hat{s})}$$

$$E[r] = R(x_t, a_t)$$

A Family of Estimators

TD(0) for V^π

- Suppose we have a **state-value function estimation** $V_\phi(s) \approx V^\pi(s)$
- Suppose we also have a **trajectory** $(s_1, a_1, r_1, \dots, s_\tau, a_\tau, r_\tau)$ generated by π where $s_{\tau+1}$ is a terminal state

The following are all valid estimators of $Q^\pi(s_1, a_1)$:

$$G_{1:1} = r_1 + \gamma V_\phi(s_2)$$

$$G_{1:2} = r_1 + \gamma r_2 + \gamma^2 V_\phi(s_3)$$

$$G_{1:\tau-1} = r_1 + \gamma r_2 + \gamma^2 r_3 + \dots + \gamma^{\tau-1} V_\phi(s_\tau)$$

$$G_{1:\tau} = r_1 + \gamma r_2 + \gamma^2 r_3 + \dots + \gamma^{\tau-1} r_\tau + \gamma^\tau V_\phi(s_{\tau+1})$$

$$G_{1:\tau+1} = r_1 + \gamma r_2 + \gamma^2 r_3 + \dots + \gamma^{\tau-1} r_\tau$$

$$G_{1:\tau+2} = \dots$$

$$\begin{aligned} \mathbb{E}[r_1 + \gamma V_\phi(s_2)] &= R(s_1, a_1) + \sum_{s'} P(s'|s_1, a_1) V_\phi(s') \\ &\approx R(s_1, a_1) + \sum_{s'} P(s'|s_1, a_1) V^\pi(s') = Q^\pi(s_1, a_1) \\ R(s_1, a_1) + \gamma R(s_2, a_2) &\rightarrow \sum_{s'} P(s|s_1, a_1) V_\phi(s') \\ &= R(s_1, a_1) + \gamma R(s_2, a_2) + \gamma^2 \sum_{s'} P(s|s_1, a_1) V^\pi(s') \\ &= Q^\pi(s_1, a_1) \end{aligned}$$

A Family of Estimators

$$Q^\pi(s_1, a_1) - V^\pi(s_1) = \text{advantage}$$

And the following are estimators of $Q^\pi(s_1, a_1) - V_\phi^\pi(s_1)$ (baseline)

$$A_{1:1} = r_1 + \gamma V_\phi(s_2) - V_\phi(s_1)$$

...

$$A_{1:\tau-1} = r_1 + \gamma r_2 + \gamma^2 r_3 + \dots + \gamma^{\tau-1} V_\phi(s_\tau) - V_\phi(s_1)$$

$$A_{1:\tau} = r_1 + \gamma r_2 + \gamma^2 r_3 + \dots + \gamma^{\tau-1} r_\tau - V_\phi(s_1)$$

$$A_{1:\tau+1} = r_1 + \gamma r_2 + \gamma^2 r_3 + \dots + \gamma^{\tau-1} r_\tau - V_\phi(s_1)$$

...

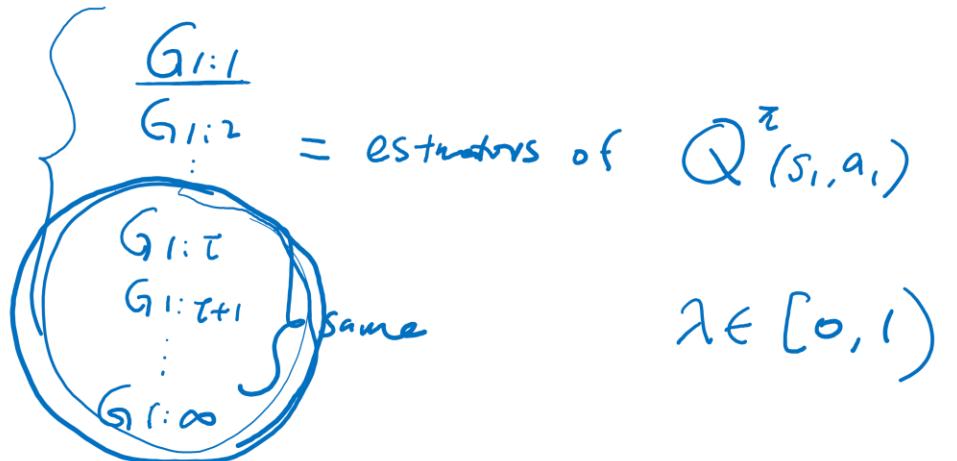
Below, we will introduce a way to combine these estimators.

Balancing Bias and Variance

$$G_1(\lambda) = (1 - \lambda) \sum_{i=1}^{\infty} \lambda^{i-1} G_{1:i}$$

$$= (1 - \lambda)(G_{1:1} + \underbrace{\lambda G_{1:2}}_{=} + \underbrace{\lambda^2 G_{1:3}}_{=} + \dots + \lambda^{\tau-1} G_{1:\tau} + \lambda^\tau G_{1:\tau+1} + \lambda^{\tau+1} G_{1:\tau+2} + \dots)$$

$$\boxed{(1-\lambda) [1 + \lambda + \lambda^2 + \dots] = 1}$$



$$\underline{A_1(\lambda)} = (1 - \lambda) \sum_{i=1}^{\infty} \lambda^{i-1} A_{1:i}$$

$$A_{1:i} = G_{1:i} - V_\phi(s_1)$$

$$\underline{\lambda \in [0, 1]}$$

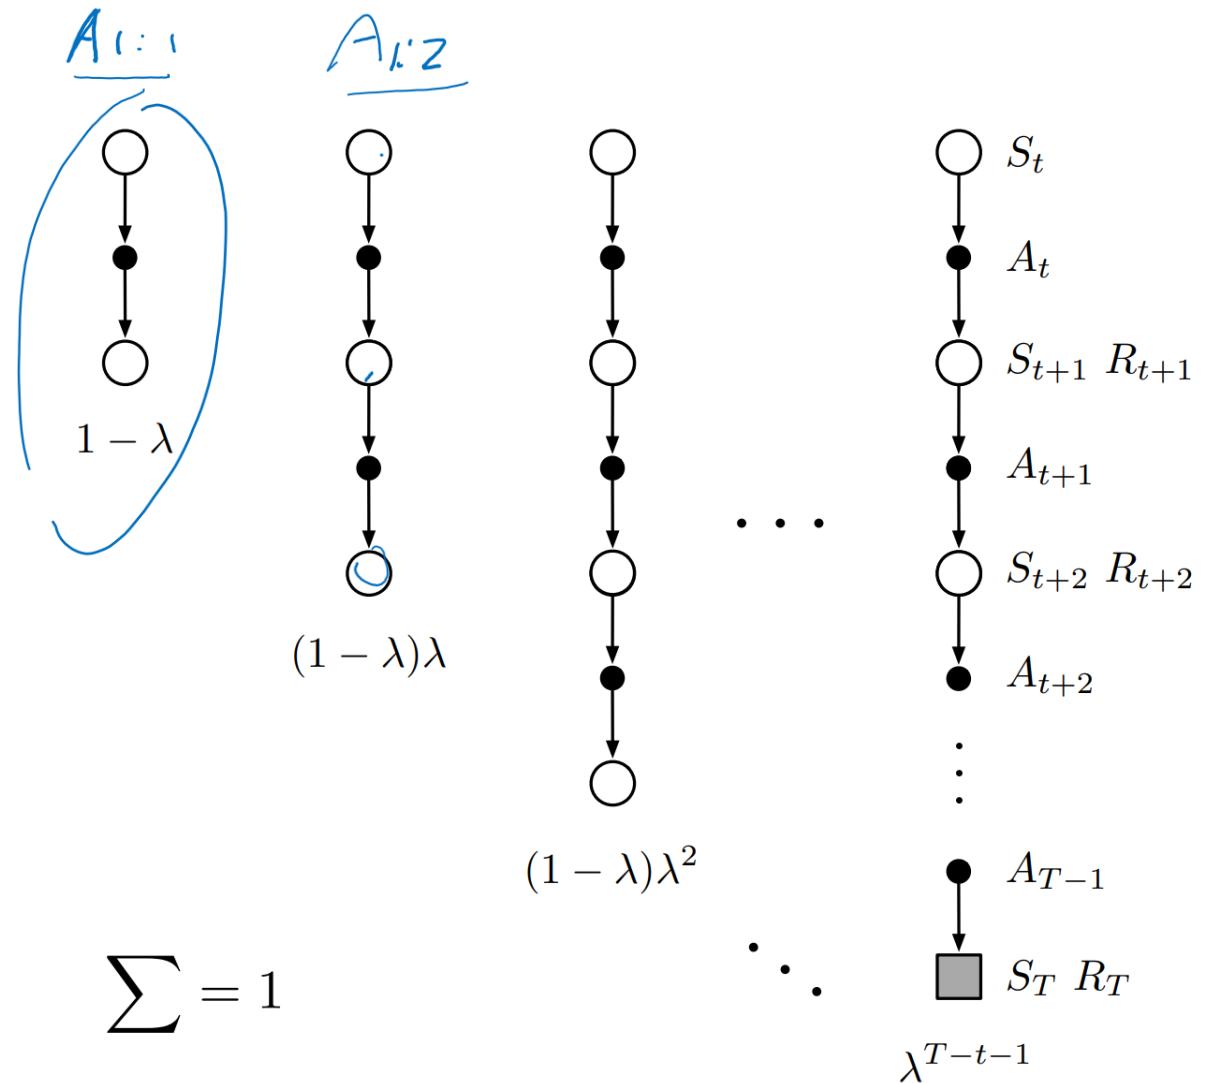
$$= (1 - \lambda)(A_{1:1} + \lambda A_{1:2} + \lambda^2 A_{1:3} + \dots + \lambda^{\tau-1} A_{1:\tau} + \lambda^\tau A_{1:\tau+1} + \lambda^{\tau+1} A_{1:\tau+2} + \dots)$$

$$A_1(\lambda) = G_1(\lambda) - V_\phi(s_1)$$

(Generalized Advantage Estimation)

An estimation of $Q^\pi(s_1, a_1) - V^\pi(s_1)$

Balancing Bias and Variance



Computing Generalized Advantage Estimator (GAE)

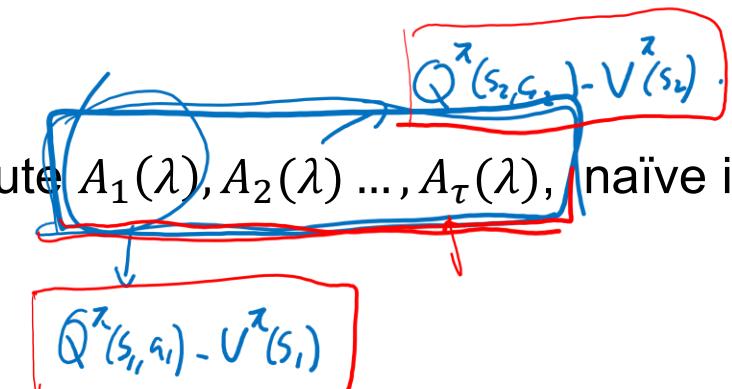
We also need to calculate

$$A_2(\lambda) = (1 - \lambda) \sum_{i=2}^{\infty} \lambda^{i-2} A_{2:i} \approx (1 - \lambda)(A_{2:2} + \lambda A_{2:3} + \lambda^2 A_{2:4} + \dots + \lambda^{\tau-2} A_{2:\tau} + \lambda^{\tau-1} A_{2:\tau+1} + \lambda^{\tau} A_{2:\tau+2} + \dots)$$

$$A_3(\lambda) = (1 - \lambda) \sum_{i=3}^{\infty} \lambda^{i-3} A_{3:i}$$

...

To compute $A_1(\lambda), A_2(\lambda), \dots, A_\tau(\lambda)$, naïve implementation takes $\tau \times \tau \times \tau$ time.



$O(\tau)$

$A_1(\lambda), \dots, A_\tau(\lambda)$

↑
Compute $A_{i:j}$ ↑
Compute $A_i(\lambda)$

Efficient Computation of GAE (1/2)

$$i \in [1, \tau]$$

Define $\delta_i = r_i + \gamma V_\phi(s_{i+1}) - V_\phi(s_i) = A_{i:i}$

$$A_{i:j} = r_i + \gamma r_{i+1} + \gamma^2 r_{i+2} + \dots + \gamma^{j-i} r_j + \gamma^{j-i+1} V_\phi(s_{j+1}) - \underline{V_\phi(s_i)}$$

$$= r_i + \gamma r_{i+1} + \gamma^2 r_{i+2} + \dots + \gamma^{j-i} r_j + \gamma^{j-i+1} V_\phi(s_{j+1}) - V_\phi(s_i)$$

$$= \delta_i + \gamma \delta_{i+1} + \gamma^2 \delta_{i+2} + \dots + \gamma^{j-i} \delta_j$$

$$\begin{aligned} &= (r_i + \gamma V_\phi(s_{i+1}) - V_\phi(s_i)) + \gamma (r_{i+1} + \gamma V_\phi(s_{i+2}) - V_\phi(s_{i+1})) + \gamma^2 (r_{i+2} + \gamma V_\phi(s_{i+3}) - V_\phi(s_{i+2})) \\ &\quad + \dots + \gamma^{j-i} (r_j + \gamma V_\phi(s_{j+1}) - V_\phi(s_j)) \end{aligned}$$

Efficient Computation of GAE (2/2)

$$A_\tau(\lambda) = (1 - \lambda)(A_{\tau:\tau} + \lambda A_{\tau:\tau+1} + \lambda^2 A_{\tau:\tau+2} + \dots) = A_{\tau\tau} = \underline{\delta_\tau} = \underline{r_\tau + \gamma V_\phi(s_{\tau+1}) - V_\phi(s_\tau)}$$

$$A_{\tau-1}(\lambda) = (1 - \lambda)(A_{\tau-1:\tau-1} + \lambda A_{\tau-1:\tau} + \lambda^2 A_{\tau-1:\tau+1} + \dots)$$

⋮

⋮

$$A_1(\lambda) = (1 - \lambda)(A_{1:1} + \lambda A_{1:2} + \lambda^2 A_{1:3} + \dots)$$

$$A_{ij} = \delta_i + \gamma \delta_{i+1} + \gamma^2 \delta_{i+2} + \dots + \gamma^{j-i} \delta_j$$

$$A_i(\lambda) = \delta_i + \lambda \gamma A_{i+1}(\lambda)$$

$$A_i(\lambda) = (1 - \lambda) \left[A_{i:i} + \lambda A_{i:i+1} + \lambda^2 A_{i:i+2} + \lambda^3 A_{i:i+3} + \dots \right]$$

$$= (1 - \lambda) \left[(\delta_i) + \lambda ((\delta_i) + \gamma \delta_{i+1}) + \lambda^2 ((\delta_i) + \gamma \delta_{i+1} + \gamma^2 \delta_{i+2}) + \lambda^3 ((\delta_i) + \gamma \delta_{i+1} + \gamma^2 \delta_{i+2} + \gamma^3 \delta_{i+3}) + \dots \right]$$

$$\approx (1 - \lambda) \left[(\underbrace{1 + \lambda + \lambda^2 + \lambda^3 + \dots}_{\text{constant}}) \delta_i + \lambda \gamma \left[(\delta_{i+1} + \lambda (\delta_{i+1} + \gamma \delta_{i+2})) + \lambda^2 (\delta_{i+1} + \lambda (\delta_{i+1} + \gamma \delta_{i+2})) + \lambda^3 (\delta_{i+1} + \lambda (\delta_{i+1} + \gamma \delta_{i+2})) + \dots \right] \right]$$

$$= \boxed{\delta_i + \lambda \gamma A_{i+1}(\lambda)}$$

GAE (Generalized Advantage Estimation)

Let $(s_1, a_1, r_1, s'_1, s_2, a_2, r_2, s'_2, \dots, s'_N, a_N, r_N, s'_N)$ be a trajectory collected with policy π , where $s'_i = s_{i+1}$ if s'_i is not a terminal state, and $s_{i+1} \sim \rho$ otherwise.

Also, let V_ϕ be a given state-value estimation.

Then the following procedure can estimate $A_i \approx \underbrace{Q^\pi(s_i, a_i) - V_\phi(s_i)}_{\text{if } s'_i \text{ is not a terminal state}} \quad i \in \{1, \dots, N\}$

Parameter: λ (controlling variance-bias tradeoff)

For $i = N, N-1, \dots, 1$:

If s'_i is a terminal state:

$$\delta_i = r_i - V_\phi(s_i)$$

$$A_i = \delta_i$$

Else:

$$\delta_i = r_i + \gamma V_\phi(s_{i+1}) - V_\phi(s_i)$$

$$A_i = \delta_i + \lambda \gamma A_{i+1}$$

Using GAE in the Policy Iteration Framework

For $k = 1, 2, \dots$

For $i = 1, 2, \dots, N$:

Choose action $a_i \sim \pi_{\theta_k}(\cdot | s_i)$

Receive reward $r_i \sim R(s_i, a_i)$ and $s'_i \sim P(\cdot | s_i, a_i)$

$s_{i+1} = s'_i$ if episode continues, $s_{i+1} \sim \rho$ if episode ends

Evaluate $Z_k(s, a) \approx Q^{\pi_{\theta_k}}(s, a) - V_{\phi}(s)$ for $s = s_1, \dots, s_N$ and all a

$$\Rightarrow Z_k(s_i, a) = \frac{\mathbb{I}\{a_i = a\}}{\pi_{\theta_k}(a|s_i)} \hat{A}_k(s_i, a_i)$$

A_{it}(x) in the previous slide

Update θ_{k+1} from θ_k using the estimator $\{Z_k(s_i, a)\}_{i=1}^N$

Using any technique we introduced for policy-based contextual bandits

Data collection

Policy Evaluation

Policy Improvement

Training the Baseline V_ϕ (in iteration k)

For $k = 1, 2, \dots$

For $i = 1, 2, \dots, N$:

Choose action $a_i \sim \pi_{\theta_k}(\cdot | s_i)$

Receive reward $r_i \sim R(s_i, a_i)$ and $s'_i \sim P(\cdot | s_i, a_i)$

$s_{i+1} = s'_i$ if episode continues, $s_{i+1} \sim \rho$ if episode ends

$$V_\phi(s) \approx V^{\pi_{\theta_k}}(s)$$

$$V^\pi(s)$$

$$r_i + \gamma V_{\phi_k}(s_{i+1})$$

$$r_i + \gamma r_{i+1} + \gamma^2 V_{\phi_k}(s_{i+2})$$

TD(0)

$$\phi_{k+1} \leftarrow \phi_k - \alpha \nabla_\phi \frac{1}{N} \sum_{i=1}^N \left(V_\phi(s_i) - r_i - \gamma V_{\phi_k}(s'_i) \right)^2 \Big|_{\phi=\phi_k}$$

$$\phi_{k+1} \leftarrow \phi_k - \alpha \nabla_\phi \frac{1}{N} \sum_{i=1}^N \left(V_\phi(s_i) - G_i(\lambda; \phi_k) \right)^2 \Big|_{\phi=\phi_k}$$

where $G_i(\lambda; \phi_k) = A_i(\lambda; \phi_k) + V_{\phi_k}(s_i)$ TD(λ)

$$\phi_{k+1} \leftarrow \phi_k - \alpha \nabla_\phi \frac{1}{N} \sum_{i=1}^N \left(V_\phi(s_i) - \sum_{h=i}^{\tau(i)} \gamma^{h-i} r_h \right)^2 \Big|_{\phi=\phi_k}$$

TD(1)

Approximate Policy Iteration and Variants

PPO

For $k = 1, 2, \dots$

For $i = 1, 2, \dots, N$:

Choose action $a_i \sim \pi_{\theta_k}(\cdot | s_i)$

Receive reward $r_i \sim R(s_i, a_i)$ and $s'_i \sim P(\cdot | s_i, a_i)$

$s_{i+1} = s'_i$ if episode continues, $s_{i+1} \sim \rho$ if episode ends

Define $Z_k(s_i, a) = \frac{\mathbb{I}\{a_i=a\}}{\pi_{\theta_k}(a|s_i)} \hat{A}_k(s_i, a_i)$

GAE $\hat{A}_k(x)$

Requires training a separate V_ϕ

Use another inner for-loop to solve the argmax with gradient ascent

$$\theta_{k+1} = \operatorname{argmax}_{\theta} \left\{ \frac{1}{N} \sum_{i=1}^N \left(\sum_a \pi_{\theta}(a|s_i) Z_k(s_i, a) - \frac{1}{\eta} \text{KL}(\pi_{\theta_k}(\cdot|s_i), \pi_{\theta}(\cdot|s_i)) \right) \right\}$$

$$\approx \operatorname{argmax}_{\theta} \left\{ \frac{1}{N} \sum_{i=1}^N \left(\frac{\pi_{\theta}(a_i|s_i)}{\pi_{\theta_k}(a_i|s_i)} \hat{A}_k(s_i, a_i) - \frac{1}{\eta} \left(\frac{\pi_{\theta}(a_i|s_i)}{\pi_{\theta_k}(a_i|s_i)} - 1 - \log \frac{\pi_{\theta}(a_i|s_i)}{\pi_{\theta_k}(a_i|s_i)} \right) \right) \right\}$$

PPO with Clipping

Schulman et al. Proximal Policy Optimization Algorithms. 2017.

$$\theta_{k+1} = \operatorname{argmax}_{\theta} \left\{ \frac{1}{N} \sum_{i=1}^N \left(\min \left\{ \frac{\pi_{\theta}(a_i|s_i)}{\pi_{\theta_k}(a_i|s_i)}, \text{clip}_{[1-\epsilon, 1+\epsilon]} \left(\frac{\pi_{\theta}(a_i|s_i)}{\pi_{\theta_k}(a_i|s_i)} \right) \right\} \hat{A}_k(s_i, a_i), -\frac{1}{\eta} \left(\frac{\pi_{\theta}(a_i|s_i)}{\pi_{\theta_k}(a_i|s_i)} - 1 - \log \frac{\pi_{\theta}(a_i|s_i)}{\pi_{\theta_k}(a_i|s_i)} \right) \right) \right\}$$

① Clipping ② KL-regularization

$L_{CLIP} = \dots$

$\epsilon \approx 0.1, 0.2$

$\hat{A}_k(s_i, a_i)$

$\pi_{\theta}(a_i|s_i)$

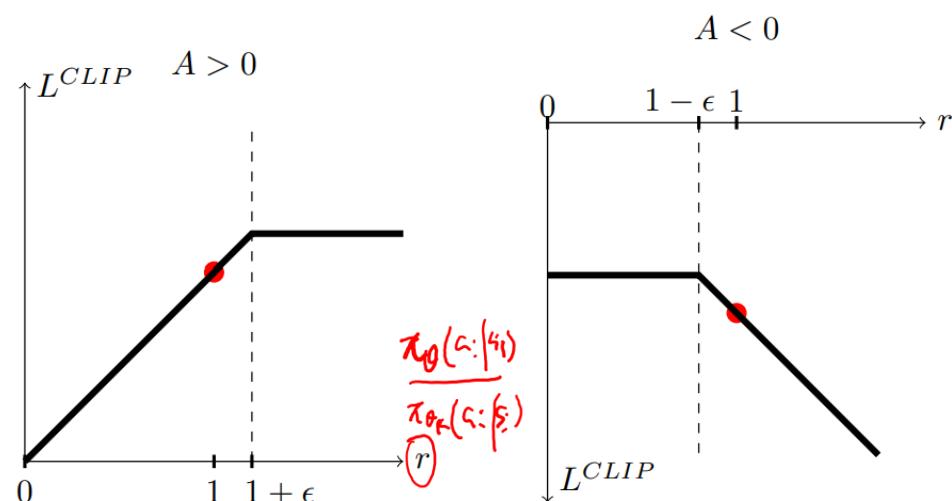
$\pi_{\theta_k}(a_i|s_i)$

$1 - \epsilon$

$1 + \epsilon$

Smaller

$1 - \epsilon$



Preventing $\frac{\pi_{\theta_{k+1}}(a_i|s_i)}{\pi_{\theta_k}(a_i|s_i)} \hat{A}_k(s_i, a_i)$ from being too high

~~from being too low~~

A2C (Advantage Actor Critic) / PG

For $k = 1, 2, \dots$

For $i = 1, 2, \dots, N$:

Choose action $a_i \sim \pi_{\theta_k}(\cdot | s_i)$

Receive reward $r_i \sim R(s_i, a_i)$ and $s'_i \sim P(\cdot | s_i, a_i)$

$s_{i+1} = s'_i$ if episode continues, $s_{i+1} \sim \rho$ if episode ends

$$\theta_{k+1} = \theta_k - \eta \frac{1}{N} \sum_{i=1}^N (\nabla_{\theta} \log \pi_{\theta}(a_i | s_i)) \Big|_{\theta=\theta_k} \hat{A}_k(s_i, a_i)$$

In standard A2C, $\hat{A}_k(s_i, a_i) = \underline{r_i + \gamma V_{\phi_k}(s'_i) - V_{\phi_k}(s_i)}$ (GAE estimator with $\lambda = 0$)

and ϕ_k is trained with TD(0):

$$\phi_{k+1} \leftarrow \phi_k - \alpha \nabla_{\phi} \frac{1}{N} \sum_{i=1}^N \left(V_{\phi}(s_i) - r_i - \gamma V_{\phi_k}(s'_i) \right)^2 \Big|_{\phi=\phi_k}$$

A2C (Advantage Actor Critic) / PG

For $k = 1, 2, \dots$

For $i = 1, 2, \dots, N$:

Choose action $a_i \sim \pi_{\theta_k}(\cdot | s_i)$

Receive reward $r_i \sim R(s_i, a_i)$ and $s'_i \sim P(\cdot | s_i, a_i)$

$s_{i+1} = s'_i$ if episode continues, $s_{i+1} \sim \rho$ if episode ends

$$\theta_{k+1} = \theta_k - \eta \frac{1}{N} \sum_{i=1}^N (\nabla_{\theta} \log \pi_{\theta}(a_i | s_i)) \Big|_{\theta=\theta_k} \hat{A}_k(s_i, a_i)$$

In standard PG, $\hat{A}_k(s_i, a_i) = \boxed{\sum_{h=i}^{\tau(i)} \gamma^{h-i} r_i} - \underline{V_{\phi_k}(s_i)}$ (GAE estimator with $\lambda = 1$)

A2C (Advantage Actor Critic) / PG

For $k = 1, 2, \dots$

For $i = 1, 2, \dots, N$:

Choose action $a_i \sim \pi_{\theta_k}(\cdot | s_i)$

Receive reward $r_i \sim R(s_i, a_i)$ and $s'_i \sim P(\cdot | s_i, a_i)$

$s_{i+1} = s'_i$ if episode continues, $s_{i+1} \sim \rho$ if episode ends

$$\theta_{k+1} = \theta_k - \eta \frac{1}{N} \sum_{i=1}^N (\nabla_{\theta} \log \pi_{\theta}(a_i | s_i)) \Big|_{\theta=\theta_k} \hat{A}_k(s_i, a_i)$$

$\cancel{V_{\phi}}, \lambda$

In general, one can use GAE with any λ to calculate $\hat{A}_k(s_i, a_i)$, with V_{ϕ} calculated from TD(λ') with any λ' .

Summary: Algorithms based on Policy Iteration

- The algorithms are almost the same as those we introduced for contextual bandits
 - PPO / NPG \leftarrow KL regularization / Clipped
 - A2C / PG \leftarrow
- The only change is replacing $r(x_i, a_i) - b(x_i)$ by Advantage Estimator:
 - $\lambda = 0$: $r(s_i, a_i) + \gamma V_\phi(s_{i+1}) - V_\phi(s_i)$
 - $\lambda = 1$: $r(s_i, a_i) + \gamma r(s_{i+1}, a_{i+1}) + \gamma^2 r(s_{i+2}, a_{i+2}) + \dots + \gamma^{\tau-i} r(s_\tau, a_\tau) - V_\phi(s_i)$
 - Any $\lambda \in [0,1]$: calculated by the GAE procedure
- The baseline $V_\phi(s)$ tries to track $V^{\pi_\theta}(s)$ where π_θ is the current policy
 - It is trained with a separate procedure TD(λ')

$$\phi_{k+1} \leftarrow \phi_k - \alpha \nabla_\phi \frac{1}{N} \sum_{i=1}^N \left(V_\phi(s_i) - r_i - \gamma V_{\phi_k}(s'_i) \right)^2 \Bigg|_{\phi=\phi_k} \quad \text{TD}(0)$$

Roadmap

