

*(for Policy Optimizer)*

# **Approximate Value Iteration and Variants**

Chen-Yu Wei

# Value Iteration

For  $k = 1, 2, \dots$

$$\forall s, a, \quad Q_k(s, a) \leftarrow \underbrace{R(s, a)}_{\text{unknown}} + \gamma \sum_{s'} \underbrace{P(s'|s, a)}_{\text{unknown}} \max_{a'} Q_{k-1}(s', a')$$

**Idea:** In each iteration, use multiple samples to estimate the right-hand side.

# Value Iteration with Samples

$$Q_{\theta_k}(x, a) \approx R(x, a)$$

For  $k = 1, 2, \dots$

$$(x_i, a_i, r_i)$$

Obtain  $N$  samples  $\{(s_i, a_i, r_i, s'_i)\}_{i=1}^N$  where  $\mathbb{E}[r_i] = R(s_i, a_i)$ ,  $s'_i \sim P(\cdot | s_i, a_i)$

Perform **regression** on  $\{(s_i, a_i, r_i, s'_i)\}_{i=1}^N$  to find  $Q_k$  such that

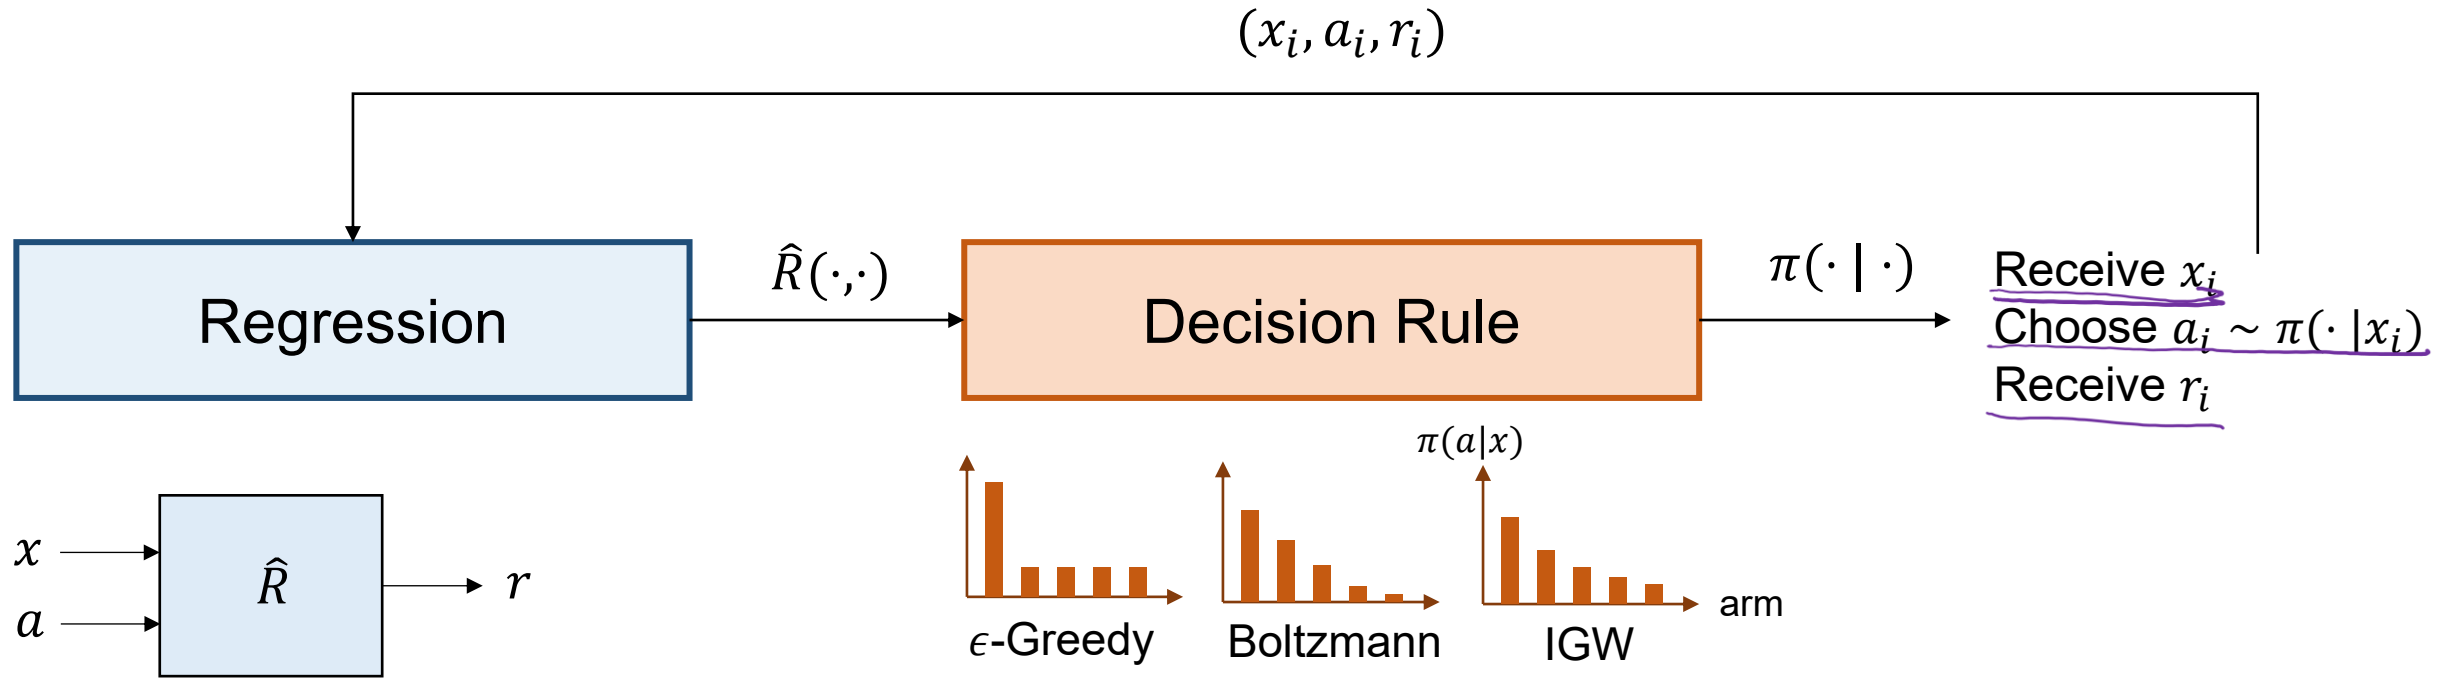
$$\forall s, a, \quad Q_k(s, a) \approx R(s, a) + \gamma \sum_{s'} P(s' | s, a) \max_{a'} Q_{k-1}(s', a')$$

Perform one iteration  
of Value Iteration

Parameterize  $Q_k = Q_{\theta_k}$

Find  $\theta_k = \underset{\theta}{\operatorname{argmin}} \sum_{i=1}^N \left( \underbrace{Q_{\theta}(s_i, a_i)} - \underbrace{r_i - \gamma \max_{a'} Q_{\theta_{k-1}}(s'_i, a')}_{} \right)^2$


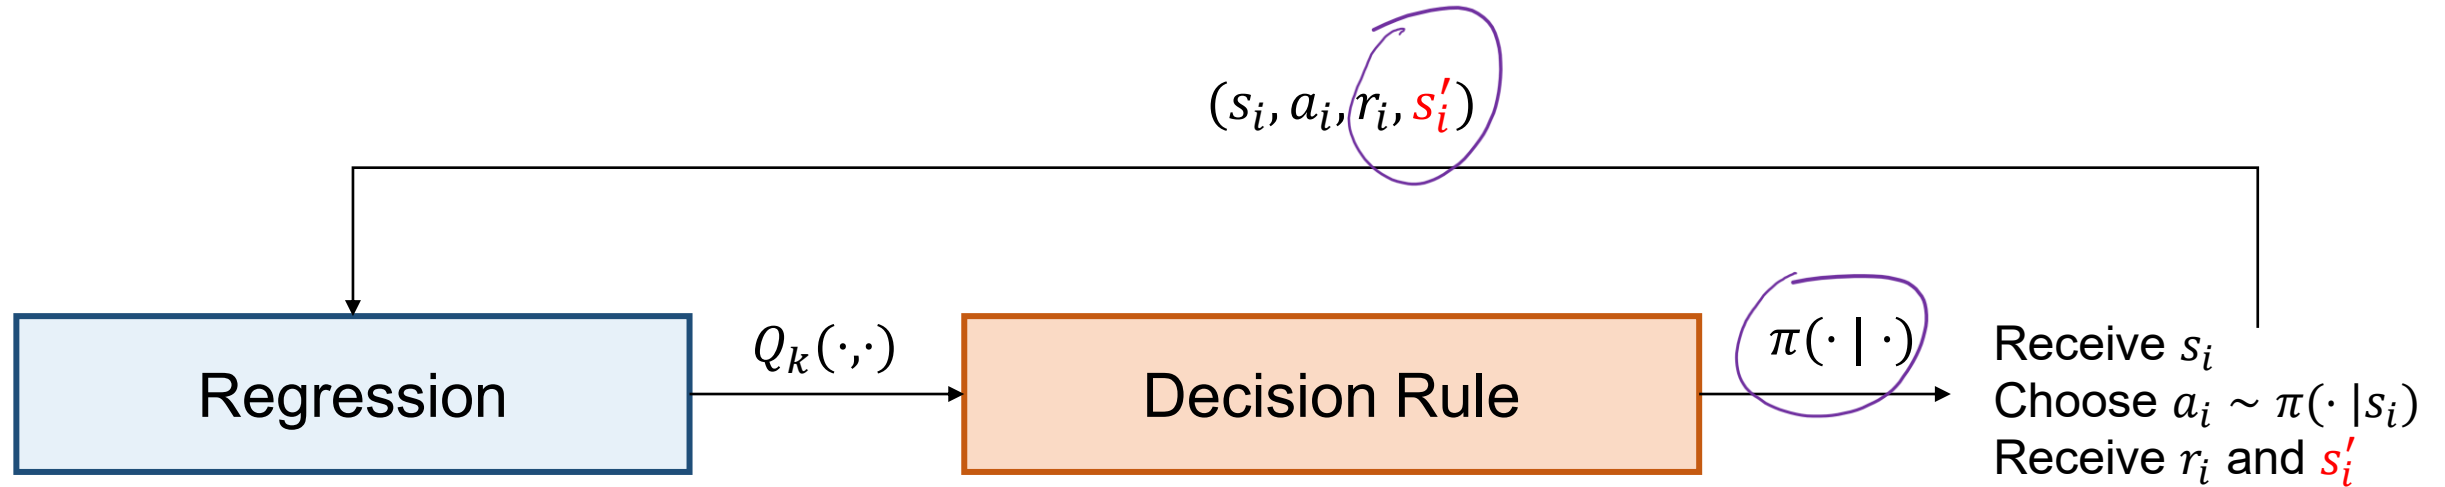
# Recall: Contextual Bandits with Regression



Train  $\hat{R}$  such that  $\hat{R}(x_i, a_i) \approx r_i = R(x_i, a_i) + \text{noise}$

$\hat{R}(x_i, a_i) \approx R(x_i, a_i)$

# Value Iteration with Regression





Train  $Q_k$  such that  $Q_k(s_i, a_i) \approx r_i + \gamma \max_{a'} Q_{k-1}(s'_i, a')$  =  $\square + \text{noise}$

This is just one iteration of Value Iteration

Handwritten notes in purple ink show the regression equation:  $Q_k(s_i, a_i) \approx \mathbb{E}_{s'_i \sim P(\cdot | s_i, a_i)} [r_i + \gamma \max_{a'} Q_{k-1}(s'_i, a')]$

# Value Iteration with Samples



For  $k = 1, 2, \dots$

For  $i = 1, 2, \dots, N$ :

Choose action  $a_i \sim \text{EG}(Q_{\theta_k}(s_i, \cdot))$  // or BE or IGW

Receive reward  $r_i \sim R(s_i, a_i)$  and  $s'_i \sim P(\cdot | s_i, a_i)$

$s_{i+1} = s'_i$  if episode continues,  $s_{i+1} \sim \rho$  if episode ends

$\theta \leftarrow \theta_k$

For  $m = 1, 2, \dots, M$ :

Randomly pick an  $i$  (or a mini-batch) from  $\{1, 2, \dots, N\}$

$\theta \leftarrow \theta - \alpha \nabla_\theta \left( Q_\theta(s_i, a_i) - r_i - \gamma \max_{a'} Q_{\theta_k}(s'_i, a') \right)^2$

$\theta_{k+1} \leftarrow \theta$

$(s_i, a_i, r_i, s'_i)$

main network  
online

target network

Data collection

$$\sum_{i=1}^N \left( Q_\theta(s_i, a_i) - Q_{\theta_k}(s_i, a_i) \right)^2$$

Perform one iteration  
of Value Iteration

Target network

**2<sup>nd</sup> for-loop:** trying to find  $\theta_{k+1} = \underset{\theta}{\operatorname{argmin}} \sum_{i=1}^N \left( Q_\theta(s_i, a_i) - r_i - \gamma \max_{a'} Q_{\theta_k}(s'_i, a') \right)^2$

# Reusing Samples



$$\forall s,a : Q_k(s,a) = R(s,a) + \gamma \mathbb{E}_{s' \sim P(\cdot|s,a)} \max_{a'} Q_{k-1}(s',a')$$

For any  $s_i, a_i$  collected in previous iters

$$Q_k(s_i, a_i) = (R(s_i, a_i)) + \gamma \mathbb{E}_{s' \sim \mathcal{D}_k} \max Q_{k-1}(\cdot)$$


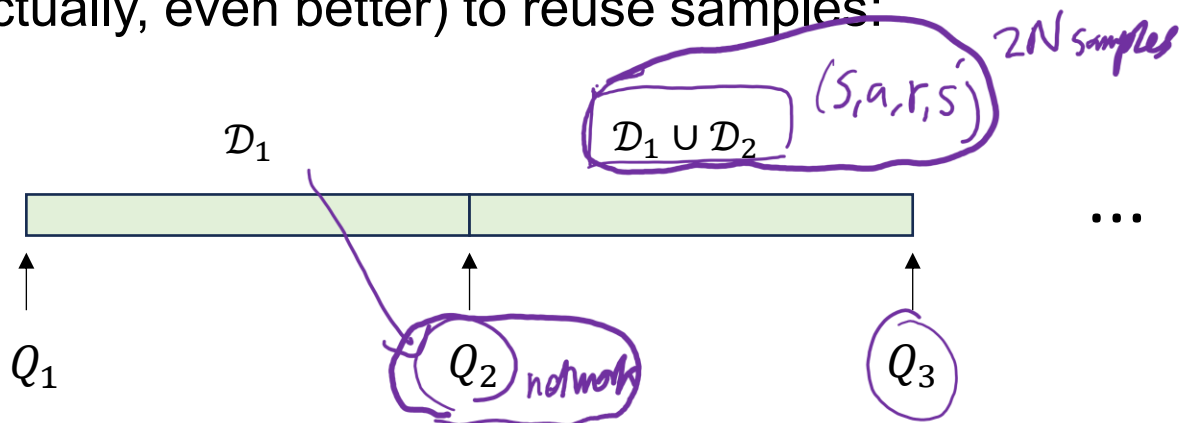
(e.g., using  $\epsilon$ -greedy)

$$\mathcal{D}_1 = \{(s_i, a_i, r_i, s'_i)\}$$



The algorithm in the previous slide only use  $\mathcal{D}_k$  to train  $\theta_{k+1}$ .

However, as the reward function  $R$  and transition  $P$  remains unchanged, it is valid (actually, even better) to reuse samples:



# Benefits of Reusing Samples

- Improving data efficiency
  - Every sample is used multiple times in training – just like we usually go through multiple epochs for supervised learning tasks.
- The ~~buffer~~ will consist of a wider range of state-actions
  - It allows <sup>Set of samples</sup> better approximation of

$$\forall s, a, \quad Q_{k+1}(s, a) = R(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_k(s', a')$$



# Value Iteration with Reused Samples (= Deep Q-Learning or DQN)

Initialize  $\mathcal{B} = \{\}$   $\leftarrow$  Replay buffer

For  $k = 1, 2, \dots$

For  $i = 1, 2, \dots, N$ :

Choose action  $a_i \sim \text{EG}(Q_{\theta_k}(s_i, \cdot))$  // or BE or IGW

Receive reward  $r_i \sim R(s_i, a_i)$  and  $s'_i \sim P(\cdot | s_i, a_i)$

$s_{i+1} = s'_i$  if episode continues,  $s_{i+1} \sim \rho$  if episode ends

Push  $(s_i, a_i, r_i, s'_i)$  to  $\mathcal{B}$

$\theta \leftarrow \theta_k$

For  $m = 1, 2, \dots, M$ :

Randomly pick an  $i$  (or a mini-batch) from  $\mathcal{B}$

$\theta \leftarrow \theta - \alpha \nabla_{\theta} \left( \underbrace{Q_{\theta}(s_i, a_i)} - r_i - \gamma \max_{a'} \underbrace{Q_{\theta_k}(s'_i, a')} \right)^2$

$\theta_{k+1} \leftarrow \theta$

HW3 task

Data collection

Perform one iteration  
of Value Iteration

Target network

# Another Popular Implementation

HW3 task

Initialize  $\mathcal{B} = \{\}$   $\leftarrow$  Replay buffer

For  $k = 1, 2, \dots$

For  $i = 1, 2, \dots, N$ :

Choose action  $a_i \sim \text{EG}(Q_\theta(s_i, \cdot))$

Receive reward  $r_i \sim R(s_i, a_i)$  and  $s'_i \sim P(\cdot | s_i, a_i)$

$s_{i+1} = s'_i$  if episode continues,  $s_{i+1} \sim \rho$  if episode ends

Push  $(s_i, a_i, r_i, s'_i)$  to  $\mathcal{B}$

For  $m = 1, 2, \dots, M$ :

Randomly pick an  $i$  (or a mini-batch) from  $\mathcal{B}$

$$\theta \leftarrow \theta - \nabla_\theta \left( Q_\theta(s_i, a_i) - r_i - \gamma \max_{a'} Q_{\bar{\theta}}(s'_i, a') \right)^2$$

$$\bar{\theta} \leftarrow (1 - \tau)\bar{\theta} + \tau\theta$$

$\bar{\theta}$   
 $\uparrow$   
Target network

$$\tau \approx 0.01$$

better approximate VI

# When Does DQN Succeed?

DQN tries to approximate **Value Iteration** by solving

$$\theta_{k+1} = \underset{\theta}{\operatorname{argmin}} \sum_{i \in \mathcal{B}} \left( Q_{\theta}(s_i, a_i) - r_i - \gamma \max_{a'} Q_{\theta_k}(s'_i, a') \right)^2 \quad (1)$$


The true Value Iteration:

$$\forall s, a, \quad Q_{k+1}(s, a) = R(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_k(s', a') \quad (2)$$

Under what conditions can (1) well approximate (2)?

- $\mathcal{B}$  should contain a wide range of state-action pairs (a challenge of **exploration**) ✓
- $Q_{\theta_{k+1}}(s, a)$  should recover  $R(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_{\theta_k}(s', a')$  well for all state-actions (a challenge of **function approximation**, or **generalization**)

# 1. Exploration in MDPs (Not Easy)



Environment:

- Fixed-horizon MDP with episode length  $H$
- Initial state at 0
- A single rewarding state at state  $H$
- Actions: Go LEFT or RIGHT


$$\frac{1}{2^H}$$

Suppose we perform DQN with  $\epsilon$ -greedy with random initialization

$\Rightarrow$  On average, we need  $2^H$  episodes to see the reward

(before that, we won't make any meaningful update and will just do random walk around state 0)

# 1. Exploration in MDPs (Not Easy)



Environment:

- Fixed-horizon MDP with episode length  $H$
- Initial state at 0
- A single rewarding state at state  $H$
- Actions: Go LEFT or RIGHT


Suppose we perform DQN with  $\epsilon$ -greedy with random initialization

⇒ On average, we need  $2^H$  episodes to see the reward

(before that, we won't make any meaningful update and will just do random walk around state 0)

# 1. Exploration in MDPs (Not Easy)

$$\text{loss} = \text{dist}(\text{Learner}, \text{goal})$$




Key issue:

- The  $\epsilon$ -greedy strategy (or BE, IGW) performs **action-space** exploration but not **state-space** exploration.
- This problem becomes more severe when the reward signal is sparse and the horizon length is **long**.
- To solve this, we usually require the exploration bonus (like UCB, TS), or a better **reward design**. (We will discuss them much later in the course)


At this point (for the discussion of DQN), we pretend that EG, BE, or IGW will lead to sufficient exploration over the **state space**.

# 1. Exploration in MDPs (Not Easy)

Classic sparse-reward environments:



Mountain Car



Montezuma's Revenge

## 2. Function Approximation

To make DQN well approximate VI, we need

$$\forall s, a \quad Q_{\theta_{k+1}}(s, a) \approx R(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_{\theta_k}(s', a')$$

### ( $\epsilon$ -approximate) Bellman Completeness

an assumption both on the MDP and the function expressiveness

$$\forall \theta', \exists \theta \quad \forall s, a, \quad \left| Q_{\theta}(s, a) - \left( R(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_{\theta'}(s', a') \right) \right| \leq \epsilon$$

This allows us to quantify the regression error in each iteration.



## 2. Function Approximation

In HW1 you have shown

$\epsilon$ -Greedy ensures

$$\text{Regret} \lesssim \epsilon T + \sqrt{\frac{AT \cdot \text{Err}}{\epsilon}}$$

Regression error

$$\text{Err} = \sum_{t=1}^T \left( \hat{R}_t(x_t, a_t) - R(x_t, a_t) \right)^2$$

In value-based contextual bandits, the requirement / assumption for function approximation is

$$\exists \theta \quad \forall x, a \quad R_{\theta}(x, a) \approx R(x, a)$$

In value-based MDPs, the requirement / assumption for function approximation is

$$\forall \theta', \exists \theta \quad \forall s, a \quad Q_{\theta}(s, a) \approx R(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_{\theta'}(s', a')$$

# Analysis of DQN assuming sufficient exploration and Bellman Completeness

Recall the analysis for the exact Value Iteration:

1. Value Iteration will terminate.

$$|Q_k(s, a) - Q_{k-1}(s, a)| \leq \epsilon \quad \forall s, a$$

2. When it terminates, it holds that

$$|Q_k(s, a) - Q^*(s, a)| \leq \frac{\epsilon}{1 - \gamma} \quad \forall s, a$$

3. When it terminates, it holds that

$$V^*(s) - V^{\hat{\pi}}(s) \leq \frac{2\epsilon}{(1 - \gamma)^2} \quad \forall s$$

where  $\hat{\pi}(s) = \operatorname{argmax}_a Q_k(s, a)$

$$\begin{aligned} & \max_{s,a} |Q_k(s, a) - Q_{k-1}(s, a)| \\ & \leq \gamma \max_{s,a} |Q_{k-1}(s, a) - Q_{k-2}(s, a)| \end{aligned}$$

$$\text{ValueError} \leq \frac{1}{1 - \gamma} \text{BellmanError}$$

$$\text{Suboptimality} \leq \frac{1}{1 - \gamma} \text{ValueError}$$

# DQN can be offline

Let  $\mathcal{B}$  consists of  $(s, a, r, s')$  tuples collected by a mixture of **arbitrary policies**.

Data collection

For  $k = 1, 2, \dots$

$\theta \leftarrow \theta_k$

For  $m = 1, 2, \dots, M$ :

Randomly pick an  $i$  (or a mini-batch) from  $\mathcal{B}$

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} \left( Q_{\theta}(s_i, a_i) - r_i - \gamma \max_{a'} Q_{\theta_k}(s'_i, a') \right)^2$$

$\theta_{k+1} \leftarrow \theta$

Perform Value Iteration

Again, its success relies on 1)  $\mathcal{B}$  contains data with sufficiently wide range of state-actions, 2) Bellman completeness.

The same theoretical analysis applies.

# **Handling the Non-Ideal Case**

# When DQN cannot well-approximate VI

In practice,

- We may not be able to collect sufficiently wide range of state-actions
- Bellman completeness may not hold


In either case, we may not have

$$\forall s, a \quad Q_{\theta_{k+1}}(s, a) \approx R(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_{\theta_k}(s', a')$$

This makes our previous analysis based on VI fails.

# When DQN cannot well-approximate VI

In this case,  $Q_{\theta_k}(s, a)$  tends to **overestimate**  $Q^*(s, a)$ , and the greedy policy  $\hat{\pi}(s) = \operatorname{argmax}_a Q_{\theta_k}(s, a)$  could be very bad.



# When DQN cannot well-approximate VI


Such “seeking the error” behavior is due to “**bootstrapping**”

- An issue only in MDP but not in bandits

To prevent overestimation, two strategies are

- Double Q-learning: decorrelating the choice of argmax action and the error of the value function
- Conservative Q-learning: being conservative


# Double DQN (v1)



$\theta_1$

$\bar{\theta}_1$

$$\text{loss} = \left( Q_{\theta_1}(s, a) - r - \gamma \max_{a'} Q_{\bar{\theta}_1}(s', a') \right)^2$$




$\theta_2$

$\bar{\theta}_2$


$$\text{loss} = \left( Q_{\theta_2}(s, a) - r - \gamma \max_{a'} Q_{\bar{\theta}_2}(s', a') \right)^2$$



# Double DQN (v1)




$\max_{a'} Q_{\bar{\theta}_1}(s', a')$



$\max_{a'} Q_{\bar{\theta}_2}(s', a')$

# Double DQN (v2)



Hado van Hasselt, Arthur Guez, David Silver. Deep Reinforcement Learning with Double Q-learning. 2015.

# Conservative Q-learning (CQL)

$$\begin{aligned}\theta_{k+1} &= \operatorname{argmin}_{\theta} \sum_{i \in \mathcal{B}} \left( Q_{\theta}(s_i, a_i) - r_i - \gamma \max_{a'} Q_{\theta_k}(s'_i, a') \right)^2 \\ &\quad + \alpha \sum_{i \in \mathcal{B}} \left( \log \left( \sum_a \exp(Q_{\theta}(s_i, a)) \right) - Q_{\theta}(s_i, a_i) \right) \\ &= \operatorname{argmin}_{\theta} \sum_{i \in \mathcal{B}} \left( Q_{\theta}(s_i, a_i) - r_i - \gamma \max_{a'} Q_{\theta_k}(s'_i, a') \right)^2 \\ &\quad + \alpha \sum_{i \in \mathcal{B}} \left( \max_{\mu} \sum_a \mu(a|s_i) Q_{\theta}(s_i, a) - Q_{\theta}(s_i, a_i) - \text{KL}(\mu(\cdot | s_i), \text{Unif}) \right)\end{aligned}$$

# Comparison

- Double-Q: make the  $\operatorname{argmax}_a Q_\theta(s, a)$  choice decoupled from  $\theta$
- Conservative-Q: mitigate the overestimation of  $\max_a Q_\theta(s_i, a)$  over  $Q_\theta(s_i, a_i)$


# Summary for DQN

- Motivation: approximating Value Iteration using **samples** and **function approximation**
- Standard elements: target network, replay buffer
- Work as desired when both of the following conditions hold:
  - The learner is able to obtain exploratory data (online or offline)
  - Neural network is sufficiently expressive: Bellman completeness
- When the conditions above do not hold
  - Tends to overestimate  $Q$  values and suggest arbitrary actions
- Solutions
  - Double Q-learning
  - Conservative Q-learning

# Improvements on DQN

- Dueling DDQN
- Prioritized replay
- Distributional DQN
- ...

Rainbow: Combining Improvements in Deep Reinforcement Learning. 2018.



## **Other Variants that Fail**

# An Unstable Variant

DQN without target network

For  $k = 1, 2, \dots$

Randomly pick an  $i$  (or a mini-batch) from  $\mathcal{B}$

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} \left( Q_{\theta}(s_i, a_i) - r_i - \gamma \max_{a'} Q_{\bar{\theta}}(s'_i, a') \right)^2$$

$$\bar{\theta} \leftarrow \theta$$

cf. DQN with target network

For  $k = 1, 2, \dots$

Randomly pick an  $i$  (or a mini-batch) from  $\mathcal{B}$

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} \left( Q_{\theta}(s_i, a_i) - r_i - \gamma \max_{a'} Q_{\bar{\theta}}(s'_i, a') \right)^2$$

$$\bar{\theta} \leftarrow (1 - \tau) \bar{\theta} + \tau \theta$$

For  $k = 1, 2, \dots$

$$\theta \leftarrow \bar{\theta}$$

For  $m = 1, \dots, M$ :

Randomly pick an  $i$  (or a mini-batch) from  $\mathcal{B}$



$$\theta \leftarrow \theta - \alpha \nabla_{\theta} \left( Q_{\theta}(s_i, a_i) - r_i - \gamma \max_{a'} Q_{\bar{\theta}}(s'_i, a') \right)^2$$

$$\bar{\theta} \leftarrow \theta$$



# An Unstable Variant

**Diverges** even when exploration assumption and Bellman completeness hold



Simplified from the “Baird’s counterexample”  
(see Sutton and Barto Section 11.2)

# The Effect of Target Network

Let  $KN = 100000$

For  $k = 1, 2, \dots, K$

$$\theta_k \leftarrow \theta$$


For  $i = 1, \dots, N$ :

Sample  $(s, a, r, s') \sim \text{Uniform} \{(s_1, a, 1, s_2), (s_2, a, 0, s_2)\}$


$$\theta \leftarrow \theta - \alpha \left( \phi(s, a)^\top \theta - r - \gamma \phi(s', a)^\top \theta_k \right) \phi(s, a)$$

$$\theta_{k+1} \leftarrow \theta$$


# The Effect of Target Network




N=1



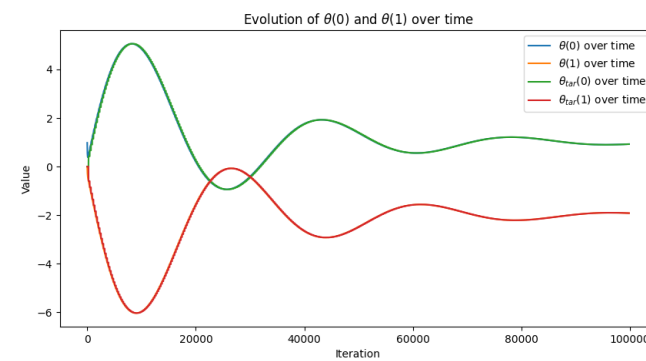
N=150




N=210




N=170




N=230




N=190




N=250




N=300




N=1000




N=500



N=2000



N=800



N=5000

# A Biased Variant

DQN without **stop gradient**

For  $k = 1, 2, \dots$

Randomly pick an  $i$  (or a mini-batch) from  $\mathcal{B}$

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} \left( Q_{\theta}(s_i, a_i) - r_i - \gamma \max_{a'} Q_{\theta}(s'_i, a') \right)^2$$

*cf.* standard DQN

For  $k = 1, 2, \dots$

Randomly pick an  $i$  (or a mini-batch) from  $\mathcal{B}$

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} \left( Q_{\theta}(s_i, a_i) - r_i - \gamma \max_{a'} Q_{\bar{\theta}}(s'_i, a') \right)^2$$

$$\bar{\theta} \leftarrow (1 - \tau) \bar{\theta} + \tau \theta$$

For  $k = 1, 2, \dots$

$$\theta \leftarrow \bar{\theta}$$

For  $m = 1, \dots, M$ :

Randomly pick an  $i$  (or a mini-batch) from  $\mathcal{B}$

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} \left( Q_{\theta}(s_i, a_i) - r_i - \gamma \max_{a'} Q_{\bar{\theta}}(s'_i, a') \right)^2$$

$$\bar{\theta} \leftarrow \theta$$

# A Biased Variant

This variant will converge (as it is similar to standard SGD), but the solution it converges to could be undesirable.

The underlying loss function of this algorithm is

$$\sum_{i \in \mathcal{B}} \left( Q_{\theta}(s_i, a_i) - r_i - \gamma \max_{a'} Q_{\theta}(s'_i, a') \right)^2$$

# Variants that Fail

- Both variants, while look somewhat reasonable, deviate from the idea of Value Iteration.