

# **Approximate Policy Iteration and Variants**

Chen-Yu Wei

# Policy Iteration

For  $k = 1, 2, \dots$

Calculate  $Q^{\pi_k}(s, a) \quad \forall s, a$

$\pi_{k+1}(s) = \operatorname{argmax}_a Q^{\pi_k}(s, a) \quad \forall s$

# Asynchronous Policy Iteration

For  $k = 1, 2, \dots$

Pick any state  $\hat{s}$

Calculate  $Q^{\pi_k}(\hat{s}, a) \quad \forall a$

$\pi_{k+1}(\hat{s}) = \operatorname{argmax}_a Q^{\pi_k}(\hat{s}, a)$

and  $\pi_{k+1}(s) = \pi_k(s) \quad \forall s \neq \hat{s}$

# Asynchronous Policy Iteration

- To improve policy, we may just evaluate  $Q^{\pi_k}$  on a particular state  $s$ .
- Of course, a **real improvement** is made only when  $\exists a$  s.t.  $Q^{\pi_k}(s, a) - V^{\pi_k}(s)$  is large.
- This is **different from Value Iteration**, where ideally, we would like to find  $Q_{k+1}$  such that  $Q_{k+1}(s, a) \approx R(s, a) + \gamma \mathbb{E}_{s' \sim P(\cdot|s, a)} \left[ \max_{a'} Q_k(s', a') \right] \quad \forall s, a$
- VI-based algorithm like DQN usually requires **stronger function approximation** that can generalize to unseen state

# Policy Iteration with Samples

For  $k = 1, 2, \dots$

For  $i = 1, 2, \dots, N$ :

Choose action  $a_i \sim \pi_{\theta_k}(\cdot | s_i)$

Receive reward  $r_i \sim R(s_i, a_i)$  and  $s'_i \sim P(\cdot | s_i, a_i)$

$s_{i+1} = s'_i$  if episode continues,  $s_{i+1} \sim \rho$  if episode ends

Data collection

Evaluate  $\hat{Q}_k(s, a) \approx Q^{\pi_{\theta_k}}(s, a)$  for  $s = s_1, \dots, s_N$  and all  $a$

or  $\hat{A}_k(s, a) \approx Q^{\pi_{\theta_k}}(s, a) - b_k(s)$  for  $s = s_1, \dots, s_N$  and all  $a$

Policy Evaluation

Update  $\theta_{k+1}$  from  $\theta_k$  using  $\hat{Q}_k(s, a)$  or  $\hat{A}_k(s, a)$

Using any technique we introduced for policy-based contextual bandits

Just replacing  $r(x, a) - b(x)$  by  $\hat{Q}_k(s, a)$  or  $\hat{A}_k(s, a)$

Policy Improvement

# **Policy Evaluation**

# Policy Evaluation

Given: a policy  $\pi$

Evaluate  $V^\pi(s)$  or  $Q^\pi(s, a)$

**On-policy policy evaluation:** the learner can execute  $\pi$  to evaluate  $\pi$

**Off-policy/offline policy evaluation:** the learner can only execute some  $\pi_b \neq \pi$ , or can only access some existing dataset to evaluate  $\pi$

## Use cases:

- Approximate policy iteration:  $\pi_k(s) = \operatorname{argmax}_a Q^{\pi_{k-1}}(s, a)$
- Estimate the value of a policy before deploying it in the real world, e.g., COVID-related border measures, economic recovery policies, or policy changes in recommendation systems.

# Value Iteration for $V^\pi / Q^\pi$

**Input:**  $\pi$

For  $k = 1, 2, \dots$

$$\forall s, \quad V_k(s) \leftarrow \sum_a \pi(a|s) \left( R(s, a) + \gamma \sum_{s'} P(s'|s, a) V_{k-1}(s') \right)$$

**Input:**  $\pi$

For  $k = 1, 2, \dots$

$$\forall s, a, \quad Q_k(s, a) \leftarrow R(s, a) + \gamma \sum_{s', a'} P(s'|s, a) \pi(a'|s') Q_{k-1}(s', a')$$



# **On-Policy Policy Evaluation**

# Temporal Difference (TD) Learning for $V^\pi$

For  $k = 1, 2, \dots$

Collect  $\{(s_i, a_i, r_i, s'_i)\}_{i=1}^N$  using policy  $\pi$

$$\theta_k \leftarrow \theta_{k-1} - \alpha \nabla_{\theta} \frac{1}{N} \sum_{i=1}^N \left( V_{\theta}(s_i) - r_i - \gamma V_{\theta_{k-1}}(s'_i) \right)^2 \Big|_{\theta = \theta_{k-1}}$$

No target network needed because this is an **on-policy** problem.

This algorithm is also called TD(0)

# Temporal Difference (TD) Learning for $Q^\pi$

For  $k = 1, 2, \dots$

Collect  $\{(s_i, a_i, r_i, s'_i)\}_{i=1}^N$  using policy  $\pi$

$$\theta_k \leftarrow \theta_{k-1} - \alpha \nabla_{\theta} \frac{1}{N} \sum_{i=1}^N \left( Q_{\theta}(s_i, a_i) - r_i - \gamma \sum_a \pi(a|s'_i) Q_{\theta_{k-1}}(s'_i, a') \right)^2 \Big|_{\theta = \theta_{k-1}}$$

No target network needed because this is an on-policy problem.

# Monte Carlo Estimation

Start from  $(s_1, a_1) = (\hat{s}, \hat{a})$  and execute policy  $\pi$  until the episode ends and obtain trajectory

$$s_1 = \hat{s}, a_1 = \hat{a}, r_1, s_2, a_2, r_2, \dots, s_\tau, a_\tau, r_\tau$$

Let  $G = \sum_{h=1}^{\tau} \gamma^{h-1} r_h$

$G$  is an unbiased estimator for  $Q^\pi(\hat{s}, \hat{a})$

**MC estimator:** unbiased, higher variance

**TD estimator:** biased, lower variance

# A Family of Estimators

Suppose we have a **state-value function estimation**  $V_\theta(s) \approx V^\pi(s)$

Suppose we also have a **trajectory**  $s_1, a_1, r_1, \dots, s_\tau, a_\tau, r_\tau$  generated by  $\pi$  where  $s_{\tau+1}$  is a terminal state

The following are all valid estimators of  $Q^\pi(s_1, a_1)$ :

$$G_{1:1} = r_1 + \gamma V_\theta(s_2)$$


...

$$G_{1:\tau-1} = r_1 + \gamma r_2 + \gamma^2 r_3 + \dots + \gamma^{\tau-1} V_\theta(s_\tau)$$

$$G_{1:\tau} = r_1 + \gamma r_2 + \gamma^2 r_3 + \dots + \gamma^{\tau-1} r_\tau$$

$$G_{1:\tau+1} = r_1 + \gamma r_2 + \gamma^2 r_3 + \dots + \gamma^{\tau-1} r_\tau$$

...



# A Family of Estimators

And the following are estimators of  $Q^\pi(s_1, a_1) - V_\theta(s_1)$  (baseline)

$$A_{1:1} = r_1 + \gamma V_\theta(s_2) - V_\theta(s_1)$$

...

$$A_{1:\tau-1} = r_1 + \gamma r_2 + \gamma^2 r_3 + \dots + \gamma^{\tau-1} V_\theta(s_\tau) - V_\theta(s_1)$$

$$A_{1:\tau} = r_1 + \gamma r_2 + \gamma^2 r_3 + \dots + \gamma^{\tau-1} r_\tau - V_\theta(s_1)$$

$$A_{1:\tau+1} = r_1 + \gamma r_2 + \gamma^2 r_3 + \dots + \gamma^{\tau-1} r_\tau - V_\theta(s_1)$$

...

Below, we will introduce a way to combine these estimators.

# Balancing Bias and Variance

$$\begin{aligned} G_1(\lambda) &= (1 - \lambda) \sum_{i=1}^{\infty} \lambda^{i-1} G_{1:i} \\ &= (1 - \lambda) (G_{1:1} + \lambda G_{1:2} + \lambda^2 G_{1:3} + \cdots + \lambda^{\tau-1} G_{1:\tau} + \lambda^{\tau} G_{1:\tau+1} + \lambda^{\tau+1} G_{1:\tau+2} + \cdots) \end{aligned}$$

$$\begin{aligned} A_1(\lambda) &= (1 - \lambda) \sum_{i=1}^{\infty} \lambda^{i-1} A_{1:i} && \textbf{(Generalized Advantage Estimation)} \\ &= (1 - \lambda) (A_{1:1} + \lambda A_{1:2} + \lambda^2 A_{1:3} + \cdots + \lambda^{\tau-1} A_{1:\tau} + \lambda^{\tau} A_{1:\tau+1} + \lambda^{\tau+1} A_{1:\tau+2} + \cdots) \end{aligned}$$

$$A_1(\lambda) = G_1(\lambda) - V_{\theta}(s_1)$$

# Computing Generalized Advantage Estimator (GAE)





# Using GAE in Policy Iteration Framework

For  $k = 1, 2, \dots$

For  $i = 1, 2, \dots, N$ :

Choose action  $a_i \sim \pi_{\theta_k}(\cdot | s_i)$

Receive reward  $r_i \sim R(s_i, a_i)$  and  $s'_i \sim P(\cdot | s_i, a_i)$

$s_{i+1} = s'_i$  if episode continues,  $s_{i+1} \sim \rho$  if episode ends

Data collection

Train  $V_\phi(s)$  using Temporal Difference Learning

Create  $\hat{A}_k(s, a) \approx Q^{\pi_{\theta_k}}(s, a) - V_\phi(s)$  for  $s = s_1, \dots, s_N$  and all  $a$

Policy Evaluation

Update  $\theta_{k+1}$  from  $\theta_k$  using  $\hat{A}_k(s, a)$

Using any technique we introduced for policy-based contextual bandits

Just replacing  $r(x, a) - b(x)$  by  $\hat{A}_k(s, a)$

Policy Improvement

# TD( $\lambda$ )

$$\text{TD}(0): \theta_{k+1} \leftarrow \theta_k - \alpha \nabla_{\theta} \left( V_{\theta}(s_1) - r_1 - \gamma V_{\theta_k}(s_2) \right)^2 \Big|_{\theta=\theta_k}$$

$$\text{TD}(\lambda): \theta_{k+1} \leftarrow \theta_k - \alpha \nabla_{\theta} \left( V_{\theta}(s_1) - G_1(\lambda) \right)^2 \Big|_{\theta=\theta_k}$$